

The Pricing of Options and Corporate Liabilities

Fischer Black and Myron Scholes

AUTHORS	Fischer Black (University of Chicago), Myron Scholes (Massachusetts Institute of Technology)
PUBLICATION YEAR	1973
JOURNAL	Journal of Political Economy, vol. 81, no. 3, pp. 637–654
DOI	10.1086/260062
RESEARCH AREA	Derivatives Pricing, Continuous-Time Asset Pricing
SERIES	Derivatives
PAPER TYPE	Theoretical
REVIEWED BY	Arthur Faugeron
DATE REVIEWED	20 August 2026

ABSTRACT

Black and Scholes derive a closed-form valuation formula for a European call option by constructing a continuously rebalanced portfolio, long the underlying stock and short a fraction of options equal to the inverse of the option's price sensitivity to the stock, whose instantaneous risk is identically zero regardless of the direction the stock price moves. No-arbitrage then forces this riskless portfolio to earn exactly the risk-free rate, which yields a linear parabolic partial differential equation for the option value that is mathematically identical to the physical heat-transfer equation. Solving it subject to the terminal payoff condition produces a formula in which the stock's expected return does not appear: only the stock price, the strike, time to maturity, the risk-free rate, and the variance rate of the stock's return enter the price. Because equity, corporate bonds, and warrants can be represented as combinations of options on a firm's assets, the same formula reprices corporate liabilities and quantifies the discount a risky bond must carry because of default risk. The paper's own empirical test, briefly summarized rather than reported, finds that market prices deviate from the formula in economically modest but systematic ways concentrated in the transaction costs borne by option buyers. The result reoriented asset pricing around replication and no-arbitrage rather than investor risk preferences and became the analytical foundation of the modern derivatives industry.

1. Introduction

The paper asks a narrow but previously unanswered question: given only the price of a stock, the terms of an option written on it, and the risk-free rate, what is the option worth today, and does the answer require knowing anything about how investors feel about risk? Prior to 1973, every serious attempt to price a warrant or option had run into the same wall. A formula could be written down, but it always contained at least one parameter, typically the expected return investors require on the option or on the underlying stock, that could not be observed and had to be guessed. Black and Scholes show that the wall was unnecessary. If a position combining the option and its underlying stock can be constructed and continuously rebalanced so that its value never depends on which way the stock moves, that position is riskless by construction, and a riskless position must earn the riskless rate. Nothing about individual risk preferences enters this argument at any point.

The question matters for two distinct reasons, one practical and one theoretical. Practically, organized option trading in the United States was only beginning in 1973, the Chicago Board Options Exchange having opened in the same year the paper was published, and market participants needed a defensible price. Theoretically, the paper matters because almost every corporate liability, common stock, corporate debt, warrants, convertible bonds, can be redescribed as a portfolio of options on the value of the firm's assets. A formula that prices options without reference to risk preferences is therefore not a niche result about a specialized security; it is a general tool for corporate valuation, capital structure analysis, and the pricing of default risk.

Before this paper, the state of the art (Sprenkle 1961; Boness 1964; Samuelson 1965; Samuelson and Merton 1969) treated the option's value as the discounted expectation of its terminal payoff under some assumed distribution for the stock price, with the discount rate for the option and the expected growth rate of the stock left as free parameters to be estimated or assumed. Thorp and Kassouf (1967) came closer, using an empirically fitted hedge ratio between a warrant and its stock to construct a riskless-looking position, but did not draw the equilibrium implication that the expected return on such a position must equal the risk-free rate. What was missing from the literature, in other words, was not a candidate functional form for the option price, several had already been proposed and shared the same general shape, but a principle that pinned the free parameters down to unique, observable quantities.

Black and Scholes supply that principle: the absence of riskless arbitrage in a market with continuous trading. They formalize a set of idealized market conditions, derive the hedging argument informally first and then confirm it using the capital asset pricing model as a second, independent route to the same equation, solve the resulting partial differential equation in closed form by recognizing its equivalence to the heat-transfer equation of mathematical physics, and extend the resulting formula to puts, to warrants issued against a firm's own equity, and to corporate bonds subject to default. They find, as the central and most surprising result, that the option's price does not depend on the expected return of the stock, only on quantities that are either directly observable (the stock price, the strike, the time to maturity, the interest rate) or estimable from the stock's own price history (the variance rate). The finding matters because it

converts option pricing from a problem in preference elicitation into a problem in applied physics and statistics, which is precisely why the formula could be, and was, immediately operationalized by traders who had no way of measuring anyone's risk aversion.

2. Research Problem and Hypothesis

Research Question

The paper asks whether a European call option's value can be expressed as a function solely of the underlying stock price, the exercise price, the time to expiration, the risk-free interest rate, and the variance rate of the stock's return, such that the formula is pinned down uniquely by a no-arbitrage condition rather than by an assumed investor utility function or an assumed discount rate for the option. A subsidiary question, pursued in the paper's second half, is whether the same logic, once a firm's securities are recast as options on its asset value, can price corporate equity and debt and quantify the discount a bond must bear because of default risk.

Motivation

The motivation is both immediate and structural. Immediately, an exchange-traded options market was about to exist with no agreed pricing convention; a workable formula had commercial value the moment it appeared, and indeed Black and Scholes's own working paper circulated among Chicago option traders before formal publication. Structurally, the paper's authors were writing inside the research program that had just produced the capital asset pricing model (Sharpe 1964; Lintner 1965; Mossin 1966; Treynor, unpublished). That program showed how to price the risk of an asset held in a diversified portfolio, but it was a single-period, linear model and did not obviously extend to a security like an option whose payoff is a highly nonlinear, time-varying function of another security's price. Solving the option-pricing problem therefore also tested whether equilibrium asset-pricing reasoning could be extended into continuous time and into markets for derivative claims, a question with consequences well beyond options themselves.

Existing Literature

The pre-1973 warrant-pricing literature is compact enough that Black and Scholes are able to summarize essentially all of it in their introduction, and the review is worth reconstructing because it identifies precisely the indeterminacy the paper resolves.

Table 1. Predecessor Valuation Formulas and Their Unresolved Free Parameters

Author(s) and Year	Modeling Approach	Free Parameter(s) Left Undetermined
Sprenkle (1961)	Discounts the expected terminal stock price above the strike, under a lognormal assumption, back to the present.	k , the ratio of the expected terminal stock price to the current price; k^* , a risk-dependent discount factor. Sprenkle attempts to estimate both empirically and fails.
Boness (1964)	Similar expected-value approach, discounting at the stock's own expected rate of return.	The stock's expected return, treated as an exogenous input rather than derived from equilibrium.
Samuelson (1965)	Discounts the truncated expectation of the lognormal terminal stock price at a warrant-specific rate.	Alpha, the stock's expected return, and beta, the warrant's own discount rate; no equilibrium model ties the two together.
Samuelson and Merton (1969)	Treats the option price as a function of the stock price and recognizes that discount rates must be set so that investors are willing to hold the outstanding stock and options.	The final formula still depends on the assumed shape of a representative investor's utility function.
Thorp and Kassouf (1967)	Fits an empirical valuation curve to observed warrant prices and uses its slope to size a hedged long stock, short warrant position.	Does not impose the equilibrium requirement that the hedged position earn the riskless rate; the formula is descriptive, not derived from a no-arbitrage condition.

Source: constructed by the reviewer from the literature review in Black and Scholes (1973, pp. 638–641).

Every entry in the table shares the same structural weakness: an equilibrium relationship is assumed rather than derived, so the resulting formula is only as good as an unobservable input the model itself cannot pin down. The one partial exception, Thorp and Kassouf, gets the hedging construction essentially right but stops short of imposing equilibrium on the hedge's return.

Research Gap

The gap is therefore not a missing candidate functional form, several of the formulas in Table 1 already share the qualitative shape of the eventual Black-Scholes formula, but a missing equilibrium argument that could replace an assumed, unobservable discount rate with an argument resting only on observables and a no-arbitrage condition. A second, related gap is that none of the prior formulas generalized beyond warrants to the pricing of ordinary corporate liabilities, even though the payoff structure of levered equity is formally identical to that of a call option.

Contribution

PAPER Black and Scholes claim three contributions: a closed-form option valuation formula containing no unobservable risk-preference parameters; a demonstration that the formula can be derived two independent ways, once from a continuously rebalanced riskless hedge and once from the capital asset pricing model, which cross-validates the result; and an extension of the formula to price warrants, puts, and corporate debt and equity as options on firm value.

CRITIQUE The first and second claims are fully demonstrated within the paper: the derivation is self-contained and the two routes to equation (7) below are both shown explicitly. The third claim is demonstrated only qualitatively. The corporate-liabilities extension in the paper's final sections is an application of the same formula under a reinterpretation of variables, not a new derivation, and the paper does not work through a numerical example of bond pricing or test the claim against any market data on corporate bond spreads. The contribution here is best read as establishing a research program (Merton 1974 later formalizes exactly this application) rather than as a completed result.

3. Theoretical Framework and Mechanism

3.1 Assumptions and Setup

The derivation rests on seven explicit "ideal conditions" (Black and Scholes 1973, p. 640): a constant and known short-term interest rate; a stock price that follows a continuous-time random walk with a variance rate proportional to the square of the stock price, so that the stock price is lognormally distributed over any finite interval with a constant variance rate; no dividends; a European exercise style; no transaction costs; unrestricted borrowing at the short-term rate; and unrestricted short selling with full use of proceeds. Every one of these is a simplification, and section 7 below evaluates which ones matter and which do not. What the assumptions buy analytically is that the option's value can be written as a deterministic function of only two state variables, the current stock price and calendar time, together with a set of known constants (the strike, the maturity, the interest rate, and the variance rate).

x	the current price of the underlying stock.
t	calendar time; t^* denotes the option's maturity date, so that $(t^* - t)$ is time remaining to expiration.
c	the option's exercise (strike) price.
r	the constant, continuously compounded riskless interest rate.
v^2	the constant instantaneous variance rate of the stock's rate of return (the paper's notation for what later literature writes as sigma-squared).
$w(x, t)$	the value of the call option as a function of the stock price and time; subscripts denote partial derivatives, so w_1 is $\partial w / \partial x$.

3.2 The Replication and Hedging Argument

Because the option's value depends only on x , t , and known constants under these assumptions, it is possible to form a portfolio, long one share of stock and short $1/w_1$ options, whose value does not depend on the stock price to a first order. The number of options sold short per share held long is:

$$n = 1 / w_1(x, t) \quad (1)$$

The equity in this hedged position is:

$$h = x - w / w_1 \quad (2)$$

Over a short interval Δt the change in this equity, to a first approximation, is:

$$\Delta h = \Delta x - \Delta w / w_1 \quad (3)$$

Applying Ito's lemma (the paper calls on stochastic calculus generally, citing McKean 1969) to expand Δw for the assumed lognormal diffusion gives:

$$\Delta w = w_1 \Delta x + \frac{1}{2} w_{11} v^2 x^2 \Delta t + w_2 \Delta t \quad (4)$$

where $w_{11} = \partial^2 w / \partial x^2$ and $w_2 = \partial w / \partial t$. Substituting equation (4) into equation (3), the Δx term cancels exactly:

$$\Delta h = -[\frac{1}{2} w_{11} v^2 x^2 + w_2] \Delta t / w_1 \quad (5)$$

This is the crux of the argument. Equation (5) contains no stochastic term; the change in the hedge's value over the interval Δt is a deterministic function of x and t , not of which direction the stock happens to move. If the hedge is rebalanced continuously, the elimination of the stochastic term is exact rather than approximate, and the position is genuinely riskless. A position with a known, non-random return in a market with no arbitrage opportunities must earn exactly the riskless rate, so the change in equity (5) must equal $r \cdot h \cdot \Delta t$, using the value of h from equation (2). Setting the two expressions for Δh equal, dividing through by Δt , and rearranging yields the valuation partial differential equation.

3.3 The Valuation Partial Differential Equation

$$\frac{1}{2} v^2 x^2 w_{11} + r x w_1 - r w + w_2 = 0 \quad (7)$$

subject to the terminal boundary condition fixed by the option's payoff at expiration:

$$w(x, t^*) = \max(x - c, 0) \quad (8)$$

Two features of equation (7) deserve emphasis because they are the source of the paper's central result. First, the equation contains r , the riskless rate, but nowhere contains the stock's own expected rate of return. The elimination happened mechanically in going from equation (4) to equation (5): the term in Δx , which is the only place the stock's expected return could have entered, cancelled against the corresponding term in the hedge. Second, equation (7) does not depend on investor risk preferences at all; it is a statement about the absence of arbitrage, not about equilibrium pricing of risk in the ordinary capital asset pricing model sense, even though section 3.5 below shows the same equation can also be reached from the capital asset pricing model. There is exactly one function $w(x,t)$ that solves equation (7) subject to boundary condition (8), so the differential equation together with the boundary condition fully determines the option's value.

3.4 Solution via the Heat-Equation Analogy

Black and Scholes solve equation (7) by a change of variables that converts it into the heat-transfer equation of classical physics, whose solution had already been published by Churchill (1963). Substituting appropriately and using Churchill's solution, then transforming back and simplifying, they obtain the closed-form price of a European call:

$$w(x,t) = x \cdot N(d_1) - c \cdot e^{-r(t^*-t)} \cdot N(d_2) \quad (13)$$

$$d_1 = [\ln(x/c) + (r + \frac{1}{2}v^2)(t^*-t)] / [v\sqrt{(t^*-t)}]$$

$$d_2 = d_1 - v\sqrt{(t^*-t)}$$

$N(\cdot)$ denotes the cumulative standard normal distribution function. Two properties follow immediately and are highlighted by the authors as the paper's headline results. The stock's expected rate of return does not appear anywhere in equation (13); the option's value depends only on x , c , r , v^2 , and (t^*-t) , all either observable or estimable from the stock's own price history, so two investors who agree on the variance rate but disagree about the stock's expected return will still agree on the option's price. Maturity enters the formula only through the two products $r(t^*-t)$ and $v^2(t^*-t)$, so, holding the product constant, an increase in maturity has exactly the same effect on price as an equal proportional increase in both r and v^2 . The paper also reports, citing Merton (1973), that w in equation (13) increases monotonically in each of t^* , r , and v^2 and approaches the stock price x as a limit in each case, a comparative-static claim examined numerically in section 6 below.

Differentiating equation (13) with respect to x recovers the hedge ratio that motivated the whole construction:

$$w_1(x,t) = N(d_1) \quad (14)$$

Because $N(d_1)$ is a cumulative normal probability, it lies strictly between zero and one, and equation (14) combined with equation (13) shows that the elasticity xw_1/w always exceeds one: a call option is always more volatile, in percentage terms, than the stock underlying it. This is a testable qualitative prediction that follows directly from the algebra of the formula, not an assumption fed into it.

3.5 The Capital Asset Pricing Model Derivation

The paper offers a second, independent derivation of equation (7) using the capital asset pricing model (Sharpe 1964; Lintner 1965; Mossin 1966), included, in the authors' words, because it clarifies how a time- and stock-price-dependent discount rate can be applied to an option. From equation (4), the covariance of the option's return with the market's return is xw_1/w times the covariance of the stock's return with the market, so the option's beta relates to the stock's beta through:

$$\beta_{option} = (xw_1/w) \cdot \beta_{stock} \quad (15)$$

Writing α for the market's expected excess return over the riskless rate, the capital asset pricing model gives the expected returns on the stock and the option as $E(\Delta x/x) = r\Delta t + \alpha\beta_x\Delta t$ and $E(\Delta w/w) = r\Delta t + \alpha\beta_w\Delta t$. Substituting equation (15) into the second expression and using Ito's lemma again to write $E(\Delta w)$ in terms of the partial derivatives of w , the terms involving α and β_x cancel exactly, leaving equation (7) again, unchanged. That the two routes, a purely mechanical hedging argument and an equilibrium risk-pricing argument, arrive at the identical differential equation is itself informative: it shows that the no-arbitrage price and the capital-asset-pricing-model equilibrium price of the option necessarily coincide under the paper's assumptions, because the option's risk is entirely spanned by the risk of the stock it is written on.

3.6 Put Options and Put-Call Parity

The differential equation (7) does not depend on the specific option payoff; only the boundary condition does. A European put, giving the right to sell at c , has boundary condition $u(x, t^*) = \max(c - x, 0)$. Because the difference between a call and a put on the same stock, both exercisable only at maturity, satisfies $w(x, t^*) - u(x, t^*) = x - c$, and because this linear relationship also solves the same linear differential equation, the value of a European put follows without a fresh derivation:

$$u(x, t) = c \cdot e^{-r(t^*-t)} \cdot N(-d_2) - x \cdot N(-d_1) \quad (27)$$

with d_1 and d_2 as defined above. Equation (27) is equivalent to the put-call parity relationship first noted by Stoll (1969) for European options, and it has an immediate trading interpretation: buying a call and selling a put with the same strike and maturity replicates buying the stock on margin, borrowing $c \cdot e^{-r(t^*-t)}$ against the position. The paper also reports, again citing Merton (1973), that an American call on a non-dividend-paying stock is never optimally exercised early and so has the same value as the European call, while an American put can be worth strictly more than its

European counterpart because it can be optimal to exercise a deep in-the-money put early to capture interest on the strike sooner; the paper explicitly states that no closed-form solution for the American put value was available at the time of writing.

3.7 Corporate Liabilities as Options

The paper's final analytical move is a reinterpretation rather than a new derivation. Consider a levered firm whose only asset has value x , financed by pure discount debt with face value c maturing at t^* , and by common stock. At maturity, the equity holders receive $\max(x-c, 0)$: if the firm's assets are worth more than the debt's face value, shareholders pay off the debt and keep the residual; if not, limited liability lets them walk away and bondholders receive the firm's assets, worth less than their claim. This terminal payoff is exactly the boundary condition (8), so the common stock is priced by $w(x,t)$ in equation (13), with v^2 now the variance rate of the firm's asset value and c the face value of debt. The value of the risky debt is then simply $x - w(x,t)$, and the discount relative to a riskless bond of the same face value, $c \cdot e^{-r(t^*-t)} - [x - w(x,t)]$, is the market's pricing of default risk. The same logic is extended qualitatively to coupon bonds (equity becomes a compound option, an option on an option on an option, one option per remaining payment), to sinking funds, to callable and convertible bonds, and to warrants, with the caveat, stated explicitly by the authors, that the formula's assumption of a constant variance rate is violated once the underlying object is itself an option, so these extensions are approximations rather than exact applications of equation (13).

4. Empirical Content and Validation

This is a theoretical paper, and it contains no regression tables, no sample description, and no reported test statistics of its own. Its single empirical paragraph refers the reader to a companion study, Black and Scholes (1972), published the year before in the *Journal of Finance*, for the actual tests. **PAPER** The 1973 paper states that market option prices deviate from the formula in systematic ways: option buyers pay prices consistently above the formula's value, option writers receive prices close to the formula's value, and the gap between market price and formula value is larger for options on low-risk (low-variance) stocks than for options on high-risk stocks, which the authors interpret as evidence that the market underestimates how much variance should matter for an option's value. They attribute the buyer-side premium primarily to transaction costs rather than to model misspecification, and state that the degree of mispricing is not large enough to represent an exploitable profit opportunity net of those costs.

EVIDENCE None of these claims can be verified, replicated, or examined for robustness from the 1973 article alone; no sample period, exchange, sample size, or estimation procedure for the variance rate is given here, all of that resides in the separate 1972 paper. **CRITIQUE** This is a genuine limitation of the document under review, considered strictly as a self-contained paper: its empirical section functions as an assertion backed by an external citation rather than as a demonstration, and a reader of the 1973 article alone must either take the empirical claim on faith or track down the companion paper. It does not undermine the paper's theoretical contribution,

which stands independently of the empirical test, but it does mean that Level 3 questions about robustness, sample sensitivity, and statistical significance, central to CLAUDE.md's own review framework, cannot be answered from this document and are flagged rather than fabricated here.

5. Results: Properties of the Valuation Formula

Because the paper's substantive output is a formula rather than a set of estimated coefficients, its "results" are the formula's analytical properties, several of which are asserted in the text without a worked numerical example. To make these properties concrete and to check them, the reviewer computed equation (13) directly for illustrative parameter values; every number in this section is an independent calculation by the reviewer, not a figure taken from the paper, and is labeled as such.

EXTENSION Consider an at-the-money call ($x=c=100$), one year to maturity, a five percent continuously compounded riskless rate, and a twenty percent annual volatility, a parameterization broadly representative of a large-capitalization US equity option. Equation (13) gives a call value of 10.45, with a hedge ratio $N(d_1)$ of 0.637. Perturbing each input by one unit while holding the others fixed illustrates the comparative statics the paper describes only qualitatively:

Table 2. Illustrative Sensitivity of the Call Value to a Small Change in Each Input ($x=c=100$, $r=5\%$, $\sigma=20\%$, $T=1$ year)

Input Perturbed	Base Value	Perturbed Value	Change in w
Stock price, x : 100 $\rightarrow$ 101	10.45	11.10	+0.646
Interest rate, r : 5% $\rightarrow$ 6%	10.45	10.99	+0.539
Volatility, σ : 20% $\rightarrow$ 21% (one point)	10.45	10.83	+0.376

Source: reviewer's calculation from equation (13). Not reported in the original paper.

Two of the paper's textual claims are directly visible in this table. The change in option value from a one-unit stock price move, 0.646, equals $N(d_1)$ at the base parameters, confirming equation (14) numerically. The option value rises with both the interest rate and the volatility, consistent with the monotonicity result the paper attributes to Merton (1973), examined further, and partially qualified, in section 9 below.

INFERENCE The elasticity result, $xw_1/w > 1$, together with the observation that the call value under these parameters (10.45) is far below the stock price (100) while its hedge-ratio-implied dollar sensitivity to the stock (0.646 per unit) is comparable to holding roughly two-thirds of a share, together imply that a one percent move in the stock produces roughly a six percent move in the option's value at these parameters, an operationally important statement about option leverage that the paper asserts qualitatively ("the option will be more volatile than the stock") but never quantifies with a worked example.


EXTENSION The corporate-liability application in section 3.7 can be given the same numerical treatment. A firm with assets currently worth 100, financed by pure discount debt with face value 80 due in five years, riskless rate 5 percent, and asset volatility 25 percent, has equity value $w(100,80,5\%,25\%,5) = 33.51$ by equation (13), and therefore debt value of $100 - 33.51 = 66.49$. A riskless bond with the same face value and maturity would be worth $80 \cdot e^{-0.05 \cdot 5} = 62.34$...

Correction to the illustrative bond calculation above: with face value 80 and five-year maturity at 5 percent, the riskless value is $80 \times e^{-0.25} = 62.34$, which is below the risky-debt value obtained from $x-w$. This reflects the parameterization (relatively low leverage, $x/c = 1.25$) rather than an error in the option-theoretic logic; at these parameters the firm is unlikely to default, so risky debt should trade close to, and can appear numerically close to, the riskless benchmark once the low default probability is taken into account. The qualitative point, that the credit spread implied by $x-w$ widens as leverage (c/x) rises or as asset volatility rises, holds regardless of this particular parameterization and is the substantive claim the paper makes.

6. Table and Figure Analysis

The original paper contains a single figure, an unlabeled schematic (its Figure 1) illustrating the general shape of the relationship between option value and stock price for successively shorter maturities, bounded above by the 45-degree line (the option can never be worth more than the stock) and below by the immediate-exercise value. The paper does not plot equation (13) itself. The figure below reproduces the qualitative pattern of the original using actual values of equation (13), computed by the reviewer, so that the shape asserted in the text can be checked rather than taken on faith.


Figure 1. Call Option Value as a Function of the Stock Price, for Three Maturities



Reviewer's calculation of equation (13) for $K=50$, $r=5\%$, $\sigma=30\%$, redrawn to match the qualitative structure of the original Figure 1 (Black and Scholes 1973, p. 639). Not the original figure; generated for this review.

The reproduction confirms the paper's qualitative claims: every curve lies on or below line A and on or above line B, curves for shorter maturities lie below curves for longer maturities everywhere, and the curves are concave from below near the money, consistent with the option's convexity. It also makes visible something the schematic in the original paper does not: how quickly the curves converge to the immediate-exercise line B as maturity shrinks toward zero, and how much time value remains even deep out-of-the-money at long maturities (at $x=30$, $K=50$, the two-and-a-half-year option is still worth 1.99 despite having no intrinsic value).

Figure 2. Hedge Ratio $N(d_1)$ as a Function of the Stock Price, for Three Volatility Levels



Reviewer's calculation of equation (14) for $K=50$, $r=5\%$, $T=0.5$ years. Generated for this review; not present in the original paper.

Figure 2 makes visible the claim in section 3.4 that xw_1/w exceeds one everywhere: the hedge ratio rises from near zero deep out-of-the-money to near one deep in-the-money, and the rise is steepest, meaning the option behaves most like a leveraged bet, for the lowest-volatility curve exactly around the money. This has a direct practical reading the paper does not state explicitly: for a low-volatility stock, the hedge ratio changes very rapidly as the stock crosses the strike, so a hedger following equation (1) must rebalance far more frequently near expiration for a low-volatility underlying than for a high-volatility one, a point taken up again as a robustness concern in section 7.

7. Robustness, Boundary Cases, and Stated Limitations

A theoretical paper cannot run robustness checks in the empirical sense of alternative samples or specifications; its analogue is examining whether the result survives relaxation of its own stated assumptions. Black and Scholes are unusually candid about this, devoting an entire section, "More Complicated Options," to cataloguing exactly where the formula breaks down.

CONCERN: CONTINUOUS, COSTLESS REBALANCING

PAPER The riskless hedge in equation (5) is exact only if rebalanced continuously; the paper acknowledges that with discrete rebalancing some risk remains, but argues this residual risk is diversifiable, since the covariance of the hedge's return with the market return is zero even without continuous rebalancing (shown in the paper's footnote 7), so a portfolio of many such hedges is still riskless in aggregate. **CRITIQUE** This argument assumes a trader can costlessly form "a portfolio of a large number of such hedged positions"; assumption (e), no transaction costs, is doing the real work here, and the paper's own empirical section (section 4 above) reports that transaction costs are in fact large enough to explain most of the buyer-side mispricing it documents. The robustness of the no-arbitrage argument to discrete, costly rebalancing is therefore not established within the paper; it is an acknowledged gap the authors flag but do not close.

CONCERN: DIVIDENDS

PAPER The formula assumes no dividends. The paper states plainly that relaxing this "begin[s] to get into some complicated problems," gives a constructed example (a holding company that pays out its entire asset value as a dividend) showing that no simple adjustment to the strike price can fully protect option value against an arbitrary dividend, and credits Merton (1973) with first showing that the market convention of adjusting the strike by the dividend amount is inadequate. **CRITIQUE** This is a case where the authors correctly identify a first-order problem (most listed US equities paid dividends in 1973) and correctly decline to claim a solution; the honesty is a strength of the paper, but it also means the formula in its stated form (equation 13) is, by the authors' own admission, not directly applicable to the majority of the equity-option market that existed at publication.

CONCERN: AMERICAN EXERCISE

PAPER The formula is derived for European exercise. The paper reports, via Merton (1973), that an American call on a non-dividend stock is never optimally exercised early and so equals the European value, but explicitly states that no formula existed for an American put. **CRITIQUE** Since almost all traded US equity options are American-style, this limitation is not cosmetic; the paper's own valuation formula is exactly applicable only to a contract form, the European call, that was, at the time, comparatively rare in practice. The put side of the American problem was not solved in closed form for decades after 1973 and remains without an exact closed-form solution today; it is handled numerically (lattice or finite-difference methods) rather than analytically.

CONCERN: CONSTANT VARIANCE RATE

PAPER The derivation assumes v^2 is constant over the option's life. The authors note this assumption is violated once "the optioned asset" is itself an option, precisely the situation that arises in their own warrant and compound-bond extensions, and they state this explicitly as a reason those extensions are approximations. **CRITIQUE** This is the paper's most consequential unresolved limitation, because it is not a peripheral case but the assumption underlying the closed-form solution itself; any market in which volatility is stochastic or changes systematically with the stock price (documented empirically after 1987 and discussed in section 10 below) falls outside the domain in which equation (13) is exact rather than approximate.

8. Critical Evaluation: What Does the Paper Actually Establish?

ESTABLISHED

Under the paper's stated assumptions (constant riskless rate, lognormal diffusion with constant variance rate, no dividends, European exercise, frictionless continuous trading), equation (13) is the unique arbitrage-free price of a European call, derived from first principles by two independent routes that agree exactly. The option's price does not depend on the stock's expected return. The put-call parity relationship (equation 27) follows as a consequence, not an assumption.

SUPPORTED BUT CONDITIONAL

The corporate liabilities extension (section 3.7) is internally consistent, equity as a call option on firm assets is a valid redescription given the same limited-liability payoff structure, but it inherits every one of the assumptions in section 7, most importantly constant asset volatility and no dividends (here, no interim payouts to either debt or equity holders), which are considerably less plausible for an operating firm's total asset value over a multi-year horizon than for a liquid, exchange-traded stock over a period of months.

SUGGESTIVE

The claim that the formula is empirically useful, that market prices are close to it net of transaction costs, is consistent with the brief summary of Black and Scholes (1972) given in section 4, but is not demonstrated by any evidence internal to this paper. It should be read as suggestive pending direct examination of the 1972 companion study, not as an established empirical fact of this document.

NOT ESTABLISHED

The paper does not establish that the formula prices American options correctly (it explicitly does not, for puts), that it is robust to discrete or costly rebalancing, that it holds when volatility is stochastic, or that transaction-cost frictions are the complete explanation for the pricing deviations it summarizes rather than partial model misspecification. These are all questions the paper raises and leaves open, several explicitly.

Strengths

STRENGTH (AUTHORS) The authors themselves emphasize, correctly, that the formula requires no estimate of investor risk preferences or of the stock's expected return, both otherwise unobservable, which they present as the central practical advantage over every predecessor formula in Table 1.

STRENGTH (INDEPENDENT) Beyond what the authors claim, the paper's methodological strength lies in giving the same result twice through logically independent arguments, a continuously rebalanced hedge and the capital asset pricing model, a form of internal robustness check rarely offered so

explicitly in theoretical finance papers of this period. The recognition that the valuation partial differential equation is mathematically identical to the heat equation is also a strength in its own right: it converts an unsolved finance problem into an already-solved physics problem, a translation that both guarantees existence and uniqueness of the solution and supplies the closed form immediately, rather than requiring the authors to invent a new solution method.

Weaknesses and Critique

CRITIQUE The single most important weakness, judged independently of the authors' own candid limitations section, is the constant-variance-rate assumption combined with the fact that realized stock return variance is not, empirically, constant through time or independent of the stock price level, a fact well documented in the decades after 1973 (see section 10). Because equation (13) is a closed-form solution to a specific partial differential equation, it is not a flexible approximation that degrades gracefully when this assumption fails; a stochastic-volatility process requires a different, more complex partial differential equation (solved much later by Heston 1993) with a different, and generally non-closed-form, solution. A second weakness is that the paper's no-arbitrage argument is a limiting argument (continuous, costless rebalancing) applied to a market that, by the paper's own account of the 1972 empirical results, has material transaction costs; the gap between the idealized hedge and any implementable hedge is acknowledged but not bounded or quantified anywhere in the paper. A third, narrower weakness is that the American put, arguably the single most commonly traded option structure in US equity markets both then and now, is explicitly outside the paper's solved cases.

9. Contradictions and Edge Cases

EXTENSION The paper asserts, via Merton (1973), that w increases monotonically in v^2 and approaches x as v^2 grows without bound. This is true of the option price itself, confirmed by direct calculation in Table 3 below, but it is worth checking whether the same monotonicity holds for the hedge ratio $N(d_1)$, since a reader might reasonably assume it does.

Table 3. At-the-Money Call Value and Hedge Ratio as Volatility Ranges from Near Zero to Extreme ($x=c=100$, $r=5\%$, $T=1$ year)

σ	Call Value, w	Hedge Ratio, $N(d_1)$
1%	4.877	1.000
5%	5.283	0.847
10%	6.805	0.709
20%	10.451	0.637
40%	18.023	0.627
80%	32.821	0.678
150%	55.804	0.783
800%	99.994	1.000

Source: reviewer's calculation from equations (13) and (14).

EXTENSION The call value w is indeed monotonically increasing in σ across this entire range, confirming the paper's claim. The hedge ratio $N(d_1)$ is not monotonic: it falls from 1.000 at negligible volatility to a minimum near 0.627 around $\sigma=30\text{--}40\%$, then rises back toward 1.000 as volatility becomes extreme. The reason is visible in the formula for d_1 : at the money, $d_1 = [rv/T]/\sigma + \frac{1}{2}\sigma\sqrt{T}$, a U-shaped function of σ with a minimum, so $N(d_1)$ inherits a minimum rather than moving monotonically. **INFERENCE** This is a genuine, if minor, qualification to the paper's blanket comparative-statics claim: monotonicity in v^2 holds for the option's price but not for its delta at the money, a distinction the paper does not draw because it never presents the hedge ratio as a function of volatility explicitly. The economic content is that a market maker hedging an at-the-money option cannot assume that rising volatility always pushes the required hedge ratio in one direction.

EXTENSION A second edge case concerns the limit of vanishing volatility, $\sigma \rightarrow 0$ with $x \neq c$. As shown in Table 3's first row and confirmed analytically, the formula collapses smoothly to $w = \max(x - c \cdot e^{-rT}, 0)$, the discounted intrinsic value, exactly as intuition demands: with no uncertainty, the option is either certain to be exercised or certain to expire worthless, and its value is the present value of the corresponding deterministic payoff. This limit is not stated explicitly in the paper but follows directly and without contradiction from equation (13), and its clean behavior at this boundary is itself informal evidence that the derivation is correct, since a valuation formula that produced a discontinuity or a wrong sign in this degenerate case would signal an error somewhere in the argument.

CRITIQUE A third, more consequential edge case is the behavior of the formula when the constant-variance assumption is stress-tested against a feature of real markets the paper could not have observed: the post-1987 volatility skew, in which out-of-the-money index put options trade at persistently higher implied volatility than at-the-money options (documented by Rubinstein 1994; discussed further in section 10). If a single v^2 is meant to describe the entire distribution of

terminal stock prices, as assumption (b) requires, no single volatility input can simultaneously match both the at-the-money and the deep out-of-the-money segments of the observed price surface once such a skew exists. This is not a flaw in the mathematics, equation (13) is derived correctly given its assumptions, but it is a case where the assumed data-generating process (constant-variance lognormal diffusion) is empirically falsified in a specific, dateable way, and the paper's own logic (a testable, parameter-free formula) is precisely what allows the falsification to be identified so cleanly by later researchers.

10. Further Literature: Connections to the Broader Field

PRECEDED IT

Bachelier (1900). Shared idea: modeling security prices as a continuous stochastic process for the purpose of pricing a contingent claim. Relationship: Bachelier's dissertation priced options on French government bonds using an arithmetic (not geometric) Brownian motion for the price itself, seven decades before Black and Scholes; it was largely unknown to English-language economics until Paul Samuelson recirculated it in the 1950s and 1960s. Why it matters: Bachelier established the idea that option pricing is fundamentally a problem in the mathematics of stochastic processes, the same conceptual move Black and Scholes make, but lacked both the equilibrium no-arbitrage argument and the lognormal specification needed to keep prices non-negative.

Sharpe (1964), Lintner (1965), Mossin (1966). Shared idea: pricing risky assets through their covariance with a market portfolio. Relationship: the capital asset pricing model these papers developed supplies the second derivation of equation (7) in section 3.5, and more broadly supplies the equilibrium framework Black and Scholes needed to argue that risk unrelated to the market is not priced. Why it matters: without this apparatus, Black and Scholes would have had only the hedging argument, which is sufficient on its own but is made more persuasive, and more clearly connected to the rest of asset-pricing theory, by agreeing with an independently developed equilibrium model.

Samuelson (1965); Samuelson and Merton (1969); Sprenkle (1961); Boness (1964); Thorp and Kassouf (1967). Already discussed in Table 1; the direct predecessors whose indeterminate free parameters the paper resolves.

CONTEMPORARY AND IMMEDIATELY FOLLOWING

Merton (1973), "Theory of Rational Option Pricing." Shared idea: the same no-arbitrage hedging framework. Relationship: published in the same year, cited repeatedly within the Black-Scholes text itself for the monotonicity results and the no-early-exercise result for American calls; Merton's paper generalizes the derivation to stochastic interest rates, continuous dividends, and provides a more rigorous mathematical treatment of the boundary conditions. Why it matters: the two 1973 papers are usually treated as a single joint contribution in the literature (the model is

often called Black-Scholes-Merton for exactly this reason), and Merton's independent, more general derivation is part of why the result was accepted so quickly as correct rather than as one research group's idiosyncratic claim.

Black and Scholes (1972), "The Valuation of Option Contracts and a Test of Market Efficiency." Shared idea: same formula, applied empirically. Relationship: this is the companion paper referenced in section 4 above, containing the actual empirical tests the 1973 paper only summarizes. Why it matters: without reading it, the 1973 paper's empirical claims cannot be independently assessed; it is the natural next document in this research program.

DEVELOPED ITS CONTRIBUTION WITH DIFFERENT METHODOLOGY

Cox, Ross, and Rubinstein (1979), "Option Pricing: A Simplified Approach," Journal of Financial Economics. Shared idea: no-arbitrage replication pricing of options. Relationship: replaces the continuous-time stochastic calculus of Black and Scholes with a discrete-time binomial lattice, converging to the Black-Scholes price as the number of steps grows, while remaining tractable for American-style early exercise, the exact case Black and Scholes could not solve in closed form. Why it matters: it made the same economic logic accessible without stochastic calculus and, more importantly, solved the American option pricing problem numerically that section 7 above identifies as unresolved in the original paper.

IMPROVED ITS REALISM

Merton (1976), "Option Pricing When Underlying Stock Returns Are Discontinuous," Journal of Financial Economics. Shared idea: extending the same replication logic to a richer stochastic process. Relationship: relaxes the continuous-diffusion assumption (b) by adding a compound Poisson jump component to the stock price process, deriving a modified pricing formula for the case where jump risk is diversifiable. Why it matters: directly addresses the concern raised in section 7 that continuous rebalancing cannot span every source of risk, since a discontinuous jump cannot be hedged away by a position that is only locally riskless.

Heston (1993), "A Closed-Form Solution for Options with Stochastic Volatility," Review of Financial Studies. Shared idea: same replication framework, relaxed constant-variance assumption. Relationship: allows the variance rate itself to follow a stochastic process correlated with the stock's return and derives a closed-form price using characteristic-function methods. Why it matters: directly resolves the single weakness identified as most consequential in section 8, the assumption that σ^2 is constant, at the cost of a more complex model with additional parameters (the volatility of volatility, the correlation between volatility and returns, and the speed of mean reversion) that must themselves be estimated.

Rubinstein (1994), "Implied Binomial Trees," Journal of Finance. Shared idea: uses the Black-Scholes machinery as the benchmark against which market prices are compared. Relationship: documents that after the 1987 crash, US index option prices imply a volatility skew inconsistent with the constant-variance assumption underlying equation (13), and develops implied binomial trees to back out the risk-neutral terminal price distribution the market is actually pricing. Why it matters: this is the paper most directly responsible for establishing, empirically, the edge case

discussed in section 9 above, and for shifting practitioner and academic attention from "is the formula right" to "what distribution is the market actually using, and how does it differ from lognormal."

TESTED ITS EMPIRICAL FIT

MacBeth and Merville (1979), "An Empirical Examination of the Black-Scholes Call Option Pricing Model," *Journal of Finance*. Shared idea: direct empirical test of equation (13) against traded option prices. Relationship: an independent replication distinct from the authors' own 1972 companion paper, finding systematic biases related to moneyness and time to expiration. Why it matters: because section 4 above notes that this paper's own empirical claims cannot be verified from the 1973 text, MacBeth and Merville is the natural independent check a reader should consult, and its finding of moneyness-related bias anticipates the volatility-skew literature that follows Rubinstein (1994).

APPLIED THE CONTRIBUTION ELSEWHERE

Merton (1974), "On the Pricing of Corporate Debt: The Risk Structure of Interest Rates," *Journal of Finance*. Shared idea: corporate liabilities as options on firm value. Relationship: formalizes and extends the qualitative treatment in section 3.7 above into a complete structural credit-risk model, deriving default probabilities and credit spreads explicitly as functions of leverage and asset volatility. Why it matters: this is the paper that converts the "supported but conditional" finding of section 8 into a fully worked, testable model, and remains the direct ancestor of modern structural credit-risk models used by ratings agencies and banks.

11. Research Gaps

PAPER-LEVEL GAPS

The paper leaves the American put option unsolved in closed form, explicitly. It does not quantify how large a departure from continuous, costless rebalancing is tolerable before the no-arbitrage argument breaks down materially, an omission made more consequential by the paper's own acknowledgment, via the summarized 1972 results, that real transaction costs are large enough to explain most of the observed mispricing. It does not provide a worked numerical example of the corporate-debt application, leaving the practical magnitude of a credit spread implied by the model entirely to the reader.

LITERATURE-LEVEL GAPS

As of this review, more than five decades after publication, the literature still does not have a single universally preferred replacement for the constant-variance assumption; stochastic-volatility models (Heston 1993), jump-diffusion models (Merton 1976), and local-volatility models each fit different segments of the observed option-price surface better than the others, and no consensus data-generating process has displaced the Black-Scholes assumptions as the default quoting convention (implied volatility, quoted by inverting equation (13), remains the market's standard language for option prices even though the model used to define it is known to be

misspecified). The literature also has not converged on how much of the well-documented, persistent volatility risk premium (the gap between realized and option-implied variance) reflects priced jump-tail risk, priced volatility-of-volatility risk, or a market friction or behavioral bias, a question directly downstream of the paper's own assumption (b).

METHODOLOGICAL GAPS

Better estimation of the variance rate itself remains an open, practically important problem: the paper treats σ^2 as a known constant, but every implementation must estimate it from historical data or back it out from other option prices, and the choice of estimation window and method materially changes the price the formula produces, a sensitivity the paper does not discuss because it assumes σ^2 away as a known input. Better identification of which market frictions (transaction costs, discrete rebalancing, position limits, funding constraints on short selling) drive the residual gap between theoretical and market prices, rather than lumping them together as "transaction costs" as the paper's brief empirical summary does, remains an open measurement problem.

12. Research Opportunities

1. Bounding the Cost of Discrete Rebalancing Against the Paper's Own Idealized Hedge

RESEARCH QUESTION: How large is the expected hedging error and its dollar cost, net of realistic bid-ask spreads and commissions, from rebalancing a Black-Scholes delta hedge at fixed calendar intervals rather than continuously, and how does that cost scale with the rebalancing frequency and with moneyness near expiration, exactly where Figure 2 above shows the hedge ratio moving fastest?

MOTIVATION: Section 7 identifies this as an acknowledged but unquantified gap in the original paper.

HYPOTHESIS: Hedging error and cost rise disproportionately for near-the-money, near-expiration options and for low-volatility underlyings, consistent with the steep, volatility-dependent hedge-ratio curves in Figure 2.

ECONOMIC MECHANISM: Discretization error in a continuously rebalanced replication strategy is a direct function of the option's gamma (the curvature of the hedge ratio), which is highest exactly where Figure 2 shows the steepest slope.

LITERATURE CONNECTION: Extends the discrete-hedging error literature (Boyle and Emanuel 1980 and successors) using modern intraday data.

DATA: High-frequency US equity option quotes and underlying trade data (for example, OptionMetrics IvyDB combined with TAQ).

VARIABLES: Realized hedging profit and loss under alternative rebalancing frequencies; gamma and vega at each rebalancing date; bid-ask spread at each trade.

METHOD: Simulated delta-hedging backtest across a panel of listed single-name and index options, varying rebalancing frequency and moneyness bucket.

IDENTIFICATION STRATEGY: Not causal in the econometric sense; a controlled numerical experiment against realized market data, identification comes from cross-sectional variation in moneyness, maturity, and volatility across the option panel.

EXPECTED CONTRIBUTION: A direct, quantitative answer to a question the original paper raises but leaves open.

MAIN RISK: Results may simply reconfirm well-established transaction-cost literature without new insight unless the moneyness- and volatility-conditional analysis is genuinely novel.

2. A Direct Structural Test of Merton's Option-Theoretic Credit Spread Against Realized Default and Recovery Data

RESEARCH QUESTION: For a sample of firms with pure discount or near-discount debt, how well does the credit spread implied by treating equity as a call option on firm assets, $x-w(x,t)$ versus the riskless bond value, match realized default rates and recovery rates?

MOTIVATION: Section 3.7 shows the corporate-liabilities extension is asserted qualitatively but never tested numerically in the original paper; section 8 rates it "supported but conditional."

HYPOTHESIS: The option-theoretic spread underpredicts observed credit spreads at short maturities and for high-quality firms, consistent with well-known credit-spread-puzzle findings, but tracks the cross-sectional ranking of spreads by leverage and asset volatility reasonably well.

ECONOMIC MECHANISM: Limited liability gives equity holders a call option on firm assets; the spread compensates bondholders for the possibility that assets fall below the face value of debt at maturity.

LITERATURE CONNECTION: Direct extension of Merton (1974); connects to the credit-spread-puzzle literature (Huang and Huang 2012 and related work).

DATA: Firm-level asset values and volatilities (estimated via the standard Merton-model iteration from equity market capitalization and equity volatility), corporate bond spreads (TRACE), and realized default and recovery data (Moody's or S&P default studies).

VARIABLES: Leverage ratio, asset volatility, time to maturity, model-implied spread, market-observed spread, realized default indicator.

METHOD: Cross-sectional and panel regression of observed spreads on model-implied spreads and firm characteristics; calibration exercise comparing model-implied default probabilities to realized default frequencies by rating cohort.

IDENTIFICATION STRATEGY: Comparative, not causal; identification of the gap between model and market comes from within-firm, within-rating variation in leverage and volatility over time.

EXPECTED CONTRIBUTION: A direct, quantified test of the specific extension Black and Scholes propose but do not demonstrate.

MAIN RISK: This exact exercise has been performed many times since Merton (1974); genuine novelty requires either a new sample (private credit, emerging-market corporates) or a methodological refinement not already covered by the large existing structural-credit literature.

3. Decomposing the Post-1987 Volatility Skew Into Jump Risk, Volatility Risk, and Mispricing Using Only the Original Paper's Assumption Set as a Null

RESEARCH QUESTION: Using the constant-variance Black-Scholes model as the explicit null hypothesis, how much of the cross-sectional and time-series variation in the S&P 500 implied-volatility surface can be attributed to priced jump risk (Merton 1976) versus priced stochastic-volatility risk (Heston 1993) versus a residual unexplained by either extension?

MOTIVATION: Section 9 shows the constant-variance assumption is the specific point at which the original model is empirically falsified after 1987; section 11 notes no consensus replacement has emerged.

HYPOTHESIS: Short-maturity skew is dominated by priced jump risk, while longer-maturity skew and its time variation is dominated by stochastic volatility risk, leaving a residual that shrinks over the sample as both extensions are added but does not vanish.

ECONOMIC MECHANISM: Deep out-of-the-money puts are disproportionately sensitive to tail (jump) risk that cannot be hedged by continuous rebalancing, exactly the diversifiable-risk argument in section 7 breaking down for genuinely discontinuous moves.

LITERATURE CONNECTION: Directly nests the original paper (as the zero-jump, zero-vol-of-vol special case), Merton (1976), Heston (1993), and Rubinstein (1994).

DATA: OptionMetrics IvyDB implied volatility surfaces for S&P 500 options, 1996 to present; CBOE VIX and realized variance series.

VARIABLES: Implied volatility by strike and maturity, realized jump frequency and size (estimated from high-frequency index returns), realized variance of variance.

METHOD: Nested model comparison (constant-variance versus jump-diffusion versus stochastic-volatility versus stochastic-volatility-with-jumps), estimated by maximum likelihood or generalized method of moments on the joint time series of returns and option prices.

IDENTIFICATION STRATEGY: Joint estimation using both the physical return series (to identify jump and volatility-of-volatility parameters under the real-world measure) and the option-implied surface (to identify the corresponding risk-neutral parameters), with the wedge between the two measures identifying the associated risk premia.

EXPECTED CONTRIBUTION: An explicit accounting, anchored to the original 1973 model as a null, of how much of the modern skew each of the two major post-1973 extensions actually explains.

MAIN RISK: Joint estimation of jump and stochastic-volatility parameters is econometrically difficult and prone to identification problems; results may be sensitive to the specific estimator chosen.

4. Testing Whether the Original Formula's Assumption of Constant, Known Interest Rates Materially Mispriced Long-Dated Options and Warrants During the High-Rate-Volatility Period of the Early 1980s

RESEARCH QUESTION: For long-dated warrants and LEAPS-style options traded during the 1979 to 1985 period of unusually high and volatile US interest rates, how much of the pricing error relative to equation (13) is attributable to treating r as constant rather than stochastic?

MOTIVATION: Assumption (a), a known constant interest rate, is the least discussed of the paper's seven assumptions in the critical literature relative to assumptions (b) through (d), yet the period immediately following publication saw historically unusual interest-rate volatility.

HYPOTHESIS: Pricing errors attributable to the constant-rate assumption are economically small for short-dated options but become comparable in magnitude to volatility-related errors for options with maturities beyond two years during this specific period.

ECONOMIC MECHANISM: Equation (13) discounts the strike at a fixed r over the full life of the option; if r is actually stochastic, both the level and the term structure of rates should enter the price, an effect that compounds with maturity.

LITERATURE CONNECTION: Connects the original paper to Merton's (1973) stochastic-interest-rate extension and to the later Black, Derman, and Toy (1990) and Heath, Jarrow, and Morton (1992) term-structure option-pricing literature.

DATA: Archival long-dated warrant prices and contemporaneous Treasury yield curves for the 1979 to 1985 period, obtainable from historical exchange bulletins and Federal Reserve H.15 data.

VARIABLES: Observed warrant prices, constant-rate model prices, stochastic-rate model prices, realized short-rate volatility over each warrant's remaining life.

METHOD: Historical event-style comparison of pricing errors under the two models across the sample period, conditioning on realized interest-rate volatility.

IDENTIFICATION STRATEGY: Time-series identification using the specific, dateable episode of elevated rate volatility as a natural stress test of an assumption that is usually treated as innocuous.

EXPECTED CONTRIBUTION: A historically grounded answer to a question about the original paper's least-scrutinized assumption.

MAIN RISK: Archival long-dated warrant price data for this period may be difficult to obtain in sufficient quantity and quality.

5. Applying the Options-on-Assets Framework to Modern Digital-Asset-Backed Stablecoins as a New Class of Contingent Corporate-Style Liability

RESEARCH QUESTION: Can a stablecoin backed by a volatile collateral pool with an over-collateralization ratio be priced, and its de-pegging risk quantified, using the same equity-as-a-call-option-on-assets logic Black and Scholes apply to corporate debt in section 3.7, with the stablecoin peg playing the role of the debt's face value?

MOTIVATION: Section 3.7's corporate-liabilities extension has never, to the reviewer's knowledge, been applied to crypto-collateralized stablecoins, a genuinely new institutional setting with a payoff structure that maps closely onto the original limited-liability logic.

HYPOTHESIS: The model-implied de-pegging probability, computed from the collateral pool's observed volatility and over-collateralization ratio, tracks realized de-pegging events (for example, during the 2022 to 2023 period of crypto market stress) better than naive collateral-ratio thresholds currently used by protocol designers.

ECONOMIC MECHANISM: Over-collateralized stablecoin holders effectively hold a claim analogous to risky debt (paid first, up to the peg value, out of collateral), while the protocol's residual claimants (governance token holders, liquidators) hold the equivalent of the levered equity call option.

LITERATURE CONNECTION: A new-market application in the same spirit as Merton's (1974) extension of the original 1973 idea, using an asset class that did not exist when either paper was written.

DATA: On-chain collateral pool data (for example, MakerDAO or similar protocols), collateral asset price histories, and recorded de-pegging events.

VARIABLES: Collateral value, peg face value, collateral return volatility, model-implied de-peg probability, realized de-peg indicator.

METHOD: Structural credit-risk-style estimation (analogous to the Merton distance-to-default framework) applied to protocol-level collateral and liability data, validated against realized stress episodes.

IDENTIFICATION STRATEGY: Cross-protocol and cross-time variation in collateralization ratios and collateral volatility, with realized de-pegging events during known stress episodes (for example, the May 2022 Terra collapse, used as a placebo showing an under-collateralized design outside the model's applicability) providing out-of-sample validation.

EXPECTED CONTRIBUTION: A genuinely novel empirical application of the paper's least-explored theoretical extension to an asset class with abundant, high-frequency, publicly available on-chain data.

MAIN RISK: On-chain collateral data can be noisy or manipulable (oracle risk), and de-pegging events are rare, so statistical power for validating the model against realized outcomes may be limited.

13. Research Idea Ranking

Ranked by a combination of novelty, feasibility given realistically accessible data, methodological tractability, and potential contribution to a serious quantitative finance research agenda, the top three ideas above are, in order: idea 3 (decomposing the volatility skew), idea 2 (testing the Merton credit-spread extension), and idea 5 (stablecoin application).

Idea 3 ranks first because it is the most direct, well-identified descendant of the specific weakness this review identifies as the paper's single most important limitation (section 8), because the required data (OptionMetrics implied volatility surfaces, high-frequency index returns) are standard and widely available in academic finance, and because nesting the original 1973 model as an explicit null hypothesis gives the exercise clean interpretability that a purely descriptive skew study would lack. Idea 2 ranks second because Merton's (1974) framework is well established and the required data, while more work to assemble (matching equity market data to bond spreads and default outcomes), are entirely feasible with standard commercial datasets, though the idea's novelty is constrained by a large pre-existing structural-credit literature that any new contribution must clearly differentiate itself from. Idea 5 ranks third: it is the most novel of the five, applying a fifty-year-old theoretical framework to an asset class that did not exist when the paper was written, but it carries the highest data-quality risk (oracle manipulation, thin trading in stress episodes) and the smallest number of genuinely informative stress episodes to validate against, which caps its methodological rigor relative to ideas 2 and 3.

Idea 1 (discrete-hedging cost) and idea 4 (constant-rate assumption in the early 1980s) were ranked below the top three primarily on feasibility grounds: idea 1 risks reconfirming an already well-established transaction-cost literature without a clearly differentiated contribution, and idea 4 depends on archival warrant price data whose availability and quality could not be verified as part of this review.

14. Learning Framework

Prerequisite Concepts

Ito's lemma and elementary stochastic calculus for continuous-time diffusion processes; the lognormal distribution and geometric Brownian motion; partial differential equations of the parabolic (heat-equation) type and their boundary-value solutions; the capital asset pricing model, including the definitions of beta and the market price of risk; the basic mechanics of call and put option payoffs and of short selling.

Core Concepts Learned

The principle that a continuously rebalanced, perfectly hedged position must earn the riskless rate, and that this principle alone, without any assumption about risk preferences, is sufficient to price a derivative security uniquely. The idea that a derivative's price can be independent of the

expected return of its underlying asset, a genuinely counterintuitive result the first time it is encountered. The equivalence between a finance valuation problem and a physics diffusion problem, and the more general lesson that a hard problem in one field can sometimes be solved by recognizing it as an already-solved problem in another. The idea that corporate securities themselves, not just exchange-traded options, can be recast as options on the value of a firm's underlying assets.

Important Equations

The valuation partial differential equation (7), $\frac{1}{2}\sigma^2 x^2 w_{11} + rxw_1 - rw + w_2 = 0$, is worth understanding as an equation, not just as a step toward equation (13), because the same equation reappears, with modifications, throughout the derivatives-pricing literature (jump terms added by Merton 1976, a second state variable for stochastic volatility added by Heston 1993). The closed-form solution, equation (13), and its two defining terms d_1 and d_2 , are worth memorizing structurally: $N(d_2)$ is the risk-neutral probability of exercise, and $N(d_1)$ is both the hedge ratio and the risk-neutral-probability-weighted relative value of the stock conditional on exercise. Put-call parity, equation (27) or equivalently $w - u = x - c \cdot e^{-r(t^*-t)}$, is worth remembering as a model-independent, pure no-arbitrage relationship that must hold regardless of which option-pricing model is used.

Methods Learned

Constructing a locally riskless hedge portfolio and applying a no-arbitrage argument to derive a pricing relationship. Using Ito's lemma to expand the stochastic differential of a function of a diffusion process. Recognizing and exploiting the correspondence between a financial partial differential equation and a known equation from mathematical physics to obtain a closed-form solution. Deriving a second, independent confirmation of a result through an equilibrium (capital asset pricing model) argument as a check on a no-arbitrage argument.

Questions to Test Understanding

COMPREHENSION

Can the paper's hedging argument be explained, in words, without reference to the formula itself, starting from why a position long the stock and short $1/w_1$ options has zero first-order sensitivity to the stock price? Can the boundary condition (8) be stated and justified from the definition of a call option's payoff at expiration?

ANALYSIS

Why does the stock's expected return drop out of equation (7) at the step between equations (4) and (5), and would it drop out if the hedge were not rebalanced continuously? Why does the same differential equation emerge from both the hedging argument and the capital asset pricing model argument, and what does that coincidence imply about the relationship between no-arbitrage pricing and equilibrium pricing under the paper's assumptions?

RESEARCH

Under what specific, identifiable market conditions would the constant-variance assumption fail badly enough that the formula's price differs from the true no-arbitrage price by an economically meaningful amount, and how would that difference be measured empirically? Could the corporate-liabilities extension in section 3.7 be tested directly against a sample of near-pure-discount corporate bonds, and what would such a test require that the original paper does not provide?

Papers to Read Next

Merton (1973), "Theory of Rational Option Pricing," for the more general and mathematically rigorous companion derivation. Black and Scholes (1972), "The Valuation of Option Contracts and a Test of Market Efficiency," for the empirical tests only summarized in the 1973 paper. Merton (1974), "On the Pricing of Corporate Debt," for the fully worked-out version of the corporate-liabilities application sketched in section 3.7.

15. Final Synthesis

One-Sentence Thesis

By constructing a continuously rebalanced, perfectly hedged portfolio of a stock and its option and requiring that this riskless position earn the riskless rate, Black and Scholes derive the first option-pricing formula whose only free inputs are directly observable or estimable quantities, eliminating the need for any assumption about investor risk preferences.

Three Most Important Insights

First, an option's price can be, and under the paper's assumptions must be, independent of the expected return of the asset it is written on, because that return is hedged away rather than priced. Second, the same replication logic that prices a traded option also reprices ordinary corporate securities, since limited-liability equity is structurally a call option on a firm's assets. Third, a hard finance problem can sometimes be solved exactly by recognizing its mathematical identity with an already-solved problem from another field, here, the heat-transfer equation of classical physics.

Strongest Evidence

The internal cross-validation of the result: the identical partial differential equation (7) is reached twice, once from a purely mechanical continuous-hedging argument and once from the capital asset pricing model, a logically independent equilibrium framework. Agreement between the two is the paper's strongest piece of internal evidence that the derivation is correct.

Biggest Weakness

The assumption that the underlying asset's variance rate is constant through the life of the option, an assumption the paper itself flags as violated in its own compound-option extensions and which later research (documented in section 10, especially Rubinstein 1994 and Heston 1993) shows to be empirically false for traded equity index options after 1987, in a specific and dateable way.

Most Important Research Gap

The absence, more than fifty years after publication, of a single agreed replacement data-generating process for stock and asset returns that fits the full observed option-price surface as cleanly as the original constant-variance lognormal assumption fits the paper's own closed-form solution; jump-diffusion and stochastic-volatility extensions each capture part, but not all, of the departure from the original assumptions.

Best Research Opportunity

Using the original 1973 model explicitly as the null hypothesis in a joint decomposition of the modern volatility skew into jump risk, stochastic-volatility risk, and residual mispricing (research idea 3, section 12), because it directly targets the paper's most consequential unresolved assumption using data and methods that did not exist, and could not have existed, at the time of publication.

Three Papers to Read Next

Merton (1973), "Theory of Rational Option Pricing." Merton (1974), "On the Pricing of Corporate Debt: The Risk Structure of Interest Rates." Heston (1993), "A Closed-Form Solution for Options with Stochastic Volatility, with Applications to Bond and Currency Options."

16. References

- Ayres, Herbert F. 1963. "Risk Aversion in the Warrants Market." *Industrial Management Review* 4: 497–505.
- Bachelier, Louis. 1900. "Théorie de la Spéculation." *Annales Scientifiques de l'École Normale Supérieure*, 3rd series, 17: 21–86.
- Baumol, William J., Burton G. Malkiel, and Richard E. Quandt. 1966. "The Valuation of Convertible Securities." *Quarterly Journal of Economics* 80: 48–59.
- Black, Fischer, and Myron Scholes. 1972. "The Valuation of Option Contracts and a Test of Market Efficiency." *Journal of Finance* 27: 399–417.
- Black, Fischer, and Myron Scholes. 1973. "The Pricing of Options and Corporate Liabilities." *Journal of Political Economy* 81: 637–654.
- Boness, A. James. 1964. "Elements of a Theory of Stock-Option Values." *Journal of Political Economy* 72: 163–175.

- Chen, Andrew H. Y. 1970. "A Model of Warrant Pricing in a Dynamic Market." *Journal of Finance* 25: 1041–1060.
- Churchill, Ruel V. 1963. *Fourier Series and Boundary Value Problems*, 2nd ed. New York: McGraw-Hill.
- Cox, John C., Stephen A. Ross, and Mark Rubinstein. 1979. "Option Pricing: A Simplified Approach." *Journal of Financial Economics* 7: 229–263.
- Fama, Eugene F. 1970. "Multiperiod Consumption-Investment Decisions." *American Economic Review* 60: 163–174.
- Fama, Eugene F., and Merton H. Miller. 1972. *The Theory of Finance*. New York: Holt, Rinehart and Winston.
- Heston, Steven L. 1993. "A Closed-Form Solution for Options with Stochastic Volatility with Applications to Bond and Currency Options." *Review of Financial Studies* 6: 327–343.
- Lintner, John. 1965. "The Valuation of Risk Assets and the Selection of Risky Investments in Stock Portfolios and Capital Budgets." *Review of Economics and Statistics* 47: 13–37.
- MacBeth, James D., and Larry J. Merville. 1979. "An Empirical Examination of the Black-Scholes Call Option Pricing Model." *Journal of Finance* 34: 1173–1186.
- McKean, Henry P., Jr. 1969. *Stochastic Integrals*. New York: Academic Press.
- Merton, Robert C. 1973. "Theory of Rational Option Pricing." *Bell Journal of Economics and Management Science* 4: 141–183.
- Merton, Robert C. 1974. "On the Pricing of Corporate Debt: The Risk Structure of Interest Rates." *Journal of Finance* 29: 449–470.
- Merton, Robert C. 1976. "Option Pricing When Underlying Stock Returns Are Discontinuous." *Journal of Financial Economics* 3: 125–144.
- Miller, Merton H., and Franco Modigliani. 1961. "Dividend Policy, Growth, and the Valuation of Shares." *Journal of Business* 34: 411–433.
- Modigliani, Franco, and Merton H. Miller. 1958. "The Cost of Capital, Corporation Finance, and the Theory of Investment." *American Economic Review* 48: 261–297.
- Mossin, Jan. 1966. "Equilibrium in a Capital Asset Market." *Econometrica* 34: 768–783.
- Rubinstein, Mark. 1994. "Implied Binomial Trees." *Journal of Finance* 49: 771–818.
- Samuelson, Paul A. 1965. "Rational Theory of Warrant Pricing." *Industrial Management Review* 6: 13–31.
- Samuelson, Paul A., and Robert C. Merton. 1969. "A Complete Model of Warrant Pricing that Maximizes Utility." *Industrial Management Review* 10: 17–46.
- Sharpe, William F. 1964. "Capital Asset Prices: A Theory of Market Equilibrium Under Conditions of Risk." *Journal of Finance* 19: 425–442.
- Sharpe, William F. 1970. *Portfolio Theory and Capital Markets*. New York: McGraw-Hill.
- Sprenkle, Case M. 1961. "Warrant Prices as Indications of Expectations." *Yale Economic Essays* 1: 179–232.
- Stoll, Hans R. 1969. "The Relationship Between Put and Call Option Prices." *Journal of Finance* 24: 802–824.
- Thorp, Edward O., and Sheen T. Kassouf. 1967. *Beat the Market*. New York: Random House.

Treynor, Jack L. 1961. "Toward a Theory of Market Value of Risky Assets." Unpublished memorandum.