

Empirical Asset Pricing via Machine Learning

Shihao Gu · Bryan T. Kelly · Dacheng Xiu

PUBLICATION	The Review of Financial Studies, Vol. 33, No. 5 (May 2020), pp. 2223–2273. Circulated as NBER Working Paper No. 25398 (December 2018, revised September 2019).
DOI	10.1093/rfs/hhaa009
RESEARCH AREA	Empirical asset pricing; statistical learning applied to the cross-section of expected stock returns
SERIES	Machine Learning & AI
PAPER TYPE	Empirical and methodological (a comparative statistical study, not a theoretical asset-pricing model)
LITERATURE REVIEW	Arthur Faugeron
DATE REVIEWED	25 September 2026

ABSTRACT

The paper treats the measurement of expected stock returns as a high-dimensional, nonparametric prediction problem and compares machine learning methods, ordinary and penalized linear regression, dimension reduction, generalized linear models, regression trees, and neural networks, on a common panel of 94 characteristics interacted with eight macroeconomic conditioning variables for essentially the full universe of CRSP common stocks from 1957 to 2016. All methods are trained, validated, and tested out of sample on a strictly forward-moving basis. The central finding is that flexible, interaction-capturing methods, trees and neural networks in particular, generate out-of-sample explanatory power for individual stock returns that is several multiples larger than the canonical linear benchmarks, and that translating these forecasts into decile portfolios produces Sharpe ratios far above a three-characteristic linear model. The gains trace specifically to the ability of nonlinear methods to capture interactions among a small, method-independent set of dominant signals concentrated in price trend, liquidity, and volatility characteristics. The authors are explicit that a statistical forecasting improvement is not itself a structural explanation, and the review evaluates that distinction closely.

01 Introduction

Gu, Kelly, and Xiu ask a question that is simple to state and expensive to answer properly: if one is agnostic about the functional form linking firm characteristics to expected returns, how much of the cross-section of stock returns is actually predictable, and which statistical method extracts that predictability most completely?

Before this paper, empirical asset pricing had accumulated a very large number of candidate predictors, characteristics, factor exposures, macro conditioning variables, almost all of them tested inside linear regressions or linear factor models. The implicit assumption in that literature is that expected returns are a linear, additive function of a handful of pre-selected variables. Two problems follow from that assumption. First, a linear model with hundreds of interacting candidate predictors is either underpowered, because degrees of freedom vanish once one begins interacting characteristics with macro states, or it must be pruned by the researcher's priors, which reopens the door to data mining and multiple-testing bias documented elsewhere in the literature (Harvey, Liu, and Zhu 2016). Second, if the true expected-return function is not linear or not additive, a linear specification is misspecified regardless of how carefully its regressors were chosen, and no amount of additional data fixes that.

The paper's response is to stop choosing the functional form in advance. It builds a single, shared feature set, characteristics interacted with macro variables, and lets nine families of statistical learning methods discover whatever structure exists in that feature space, subject only to an honest, chronologically ordered train, validate, and test split. What it finds is that the choice of functional form is not a technical detail: methods that can represent interactions among predictors (regression trees, gradient boosting, neural networks) recover several times more out-of-sample predictive power than methods that cannot (ordinary least squares, the standard three-characteristic linear model of the empirical asset-pricing literature), even though every method sees exactly the same information. The paper further shows that this predictive power converts into economically large portfolio Sharpe ratios, and that the dominant predictive signals recovered are consistent across methods, price trend, liquidity, and volatility, which is itself evidence that the exercise is recovering real structure rather than overfitting to noise. The result matters because it reframes a methodological choice, linear versus nonlinear, additive versus interactive, as the first-order determinant of how much of the cross-section of returns can be measured at all.

02 Research Problem and Hypothesis

2.1 Research Question

The paper does not ask, in the vague sense, "can machine learning predict stock returns." It asks a precise, falsifiable question: given a common information set, a fixed panel of 94 firm characteristics interacted with eight macroeconomic state variables, does the functional form

imposed on the map from that information set to next-month excess return materially change the recoverable out-of-sample explanatory power and the resulting portfolio performance, once model comparison is done under an identical, forward-only training and validation protocol. A tightly related second question follows directly: among the flexible methods, which specific feature, dimension reduction, sparsity, additive nonlinearity, or the capacity to model interactions among predictors, is responsible for whatever gain is found.

2.2 Motivation

The question matters for two audiences at once. For asset-pricing theorists, the size of the predictable component of returns, and the variables it loads on, constrains any equilibrium model of risk and expected return: a theory that cannot generate interaction effects between, say, size and momentum is arguing against evidence, not merely simplifying it. For practitioners and for the profession's own methodological hygiene, the paper matters because return prediction research has historically been conducted one linear specification at a time, which makes it very difficult to know whether a reported anomaly reflects genuine predictability or a fortunate draw among the many specifications a researcher could have tried (Harvey, Liu, and Zhu 2016; McLean and Pontiff 2016). A single, pre-registered comparison across a large family of methods, evaluated strictly out of sample, is a partial remedy to that problem, and the paper is explicit that this discipline, not any single method's cleverness, is doing much of the work.

2.3 Existing Literature

Three literatures precede the paper directly. The first is aggregate return prediction with a small number of macro predictors, most influentially Welch and Goyal (2008), who show that the standard predictive-regression variables for the equity premium are unstable and mostly fail out of sample. The second is cross-sectional characteristic-based prediction, running from the three-factor and momentum literature through Fama and French's (2008) systematic dissection of anomalies to Lewellen's (2015) demonstration that a simple three-characteristic cross-sectional regression, size, book-to-market, and momentum, already captures much of what more elaborate linear specifications capture. The third is the initial application of statistical learning tools to return prediction: lasso-based aggregate return forecasting across countries (Rapach, Strauss, and Zhou 2013), tree-based conditional portfolio sorts (Moritz and Zimmermann 2016), and neural networks applied narrowly to derivatives pricing or mortgage prepayment rather than to the broad cross-section (a lineage the paper traces back to Hutchinson, Lo, and Poggio 1994). Green, Hand, and Zhang (2017) supply the immediate empirical antecedent this paper builds on directly: a systematic re-examination of more than 300 previously published characteristics, from which a stable set of roughly 94 survives, and which the present paper adopts as its baseline predictor set rather than re-deriving it.

2.4 Research Gap

What none of these three literatures had done was hold the information set fixed and vary only the statistical mapping from that information set to expected returns, across a wide enough family of methods, and under an out-of-sample protocol strict enough to be informative about genuine

predictability rather than in-sample fit. Individual machine learning applications existed, but nothing had isolated functional form itself, linear versus nonlinear, additive versus interactive, as the object of study while controlling for the underlying information.

2.5 Contribution

The paper's stated contribution is threefold: a common, controlled comparative framework spanning linear, penalized linear, dimension-reduction, generalized additive, tree-based, and neural-network methods; an empirical finding that nonlinear, interaction-capable methods dominate, with the dominance traceable to interactions rather than to nonlinearity in general; and a demonstration that the predictive gains are economically, not merely statistically, significant once converted into portfolio returns. **CRITIQUE.** The first contribution is genuinely delivered, the comparative protocol is the paper's most durable methodological artifact and has been reused by a large subsequent literature. The second contribution is well supported for the specific interaction-versus-additive contrast the authors emphasize (Section 5 below examines this directly), though it is worth noting that isolating "interactions" as the single active ingredient still rests on comparing method families that differ along more than one dimension at a time, tree depth, regularization scheme, and optimization procedure all differ across the methods being compared, so the attribution is convincing but not fully clean. The third contribution, economic significance, is demonstrated but is also the point at which the review is most cautious, since transaction costs and capacity constraints are not fully modeled (Section 7).

03 Theoretical Framework and Statistical Mechanism

The paper is deliberately atheoretical about the economic source of predictability; its "theory" is a statistical decomposition of returns together with a menu of estimators for the conditional expectation function. Reconstructing that decomposition and each estimator precisely is necessary to understand what the empirical comparison in Sections 4 and 5 is actually testing.

3.1 The General Prediction Problem

Excess returns for stock i between month t and $t+1$ are decomposed as

$$r_{i,t+1} = E_t(r_{i,t+1}) + \varepsilon_{i,t+1} \quad E_t(r_{i,t+1}) = g^*(z_{i,t}) \quad (1)-(2)$$

$z_{i,t}$	the information set available at time t for stock i , formed by interacting stock-level characteristics with macroeconomic state variables.
$g^*(\cdot)$	the true, unknown conditional expectation function. It is assumed common across stocks and time, so that all cross-sectional and time-series variation in expected returns flows through $z_{i,t}$ rather than through stock- or period-specific parameters.
$\varepsilon_{i,t+1}$	the unforecastable residual, assumed mean-independent of $z_{i,t}$.

Every method in the paper is an estimator $g(\cdot; \theta)$ of the single unknown function $g^*(\cdot)$. This is the crucial modeling choice: nothing in equation (2) requires g to be linear, so the entire empirical exercise reduces to asking how well different function classes approximate an unknown g^* . The feature vector itself is built multiplicatively,

$$z_{i,t} = x_t \otimes c_{i,t} \quad (3)$$

where $c_{i,t}$ stacks the 94 firm-level characteristics (61 updated annually, 13 quarterly, 20 monthly, following Green, Hand, and Zhang 2017) and x_t stacks a constant with eight macroeconomic predictors taken from Welch and Goyal (2008): the dividend-price ratio, earnings-price ratio, book-to-market ratio, net equity expansion, the Treasury-bill rate, the term spread, the default spread, and stock market variance. The Kronecker interaction between the two, plus 74 industry dummies, produces 920 baseline features. **INFERENCE.** This construction already embeds an economic hypothesis that is easy to overlook: that macro conditions modulate the pricing of characteristics multiplicatively and uniformly across characteristics, rather than characteristic-by-characteristic in some more general way. The paper's own flexible methods can in principle re-discover richer interactions on top of this structure, but the multiplicative macro-characteristic interaction itself is imposed, not learned.

3.2 Estimation Objective and Loss Functions

Parameters are chosen to minimize an average loss over the pooled panel,

$$L(\theta) = (1/NT) \sum_{i,t} (r_{i,t+1} - g(z_{i,t}; \theta))^2 \quad (4)$$

For methods sensitive to the fat tails characteristic of individual stock returns, the paper substitutes a Huber loss for several estimators (elastic net, generalized linear model, boosted trees),

$$H(x; \xi) = x^2 \text{ if } |x| \leq \xi; \quad 2\xi|x| - \xi^2 \text{ if } |x| > \xi \quad (5)$$

which behaves like squared-error loss near zero but grows only linearly for large residuals, limiting the influence of extreme return observations on the fitted function. **PAPER.** The authors report that the Huber-loss variants outperform their squared-loss counterparts precisely because monthly stock returns are heavy-tailed, so a handful of extreme observations would otherwise dominate the objective.

3.3 Linear and Penalized Linear Methods

Ordinary least squares on the full 920-feature set, and on Lewellen's (2015) three-characteristic subset (size, book-to-market, momentum, denoted OLS-3), serve as the linear benchmark. Between these, the elastic net penalizes the least-squares objective with

$$\varphi(\theta; \lambda, \rho) = \lambda(1-\rho) \sum_j |\theta_j| + \frac{1}{2} \lambda \rho \sum_j \theta_j^2 \quad (6)$$

nesting the lasso ($\rho=0$) and ridge regression ($\rho=1$) as special cases; λ controls overall shrinkage intensity and is chosen on the validation sample. Elastic net remains, however, an additive, linear-in-parameters model: it can select and shrink among the 920 features but cannot represent an interaction that is not already handed to it as an explicit product term.

3.4 Dimension Reduction: PCR and PLS

Principal components regression first extracts orthogonal linear combinations of the predictors that maximize their own variance,

$$w_j = \operatorname{argmax}_w \operatorname{Var}(Zw), \text{ s.t. } w'w = 1, \operatorname{Cov}(Zw, Zw_l) = 0, l = 1, \dots, j-1 \quad (7)$$

and then regresses returns on the leading components, discarding the target variable entirely at the extraction stage. Partial least squares instead extracts components that maximize covariance with the return itself,

$$w_j = \operatorname{argmax}_w \operatorname{Cov}^2(R, Zw), \text{ s.t. } w'w = 1, \operatorname{Cov}(Zw, Zw_l) = 0 \quad (8)$$

INFERENCE. Because PLS conditions the very construction of its components on the outcome, it should dominate PCR whenever a small number of predictive linear combinations exists but is not aligned with the directions of largest total predictor variance; the paper's results (Section 5) are broadly consistent with this, though the margin between the two is modest, which itself suggests that the informative linear combinations in this data are not badly misaligned with the predictors' principal directions of variation.

3.5 Generalized Linear Model with Group Lasso

To introduce nonlinearity while remaining additive, the paper expands each predictor into a spline basis and penalizes at the group level,

$$g(z; \theta, p(\cdot)) = \sum_{j=1}^P p(z_j) \theta_j, \quad p(\cdot) = (1, z, (z-c_1)^2, \dots, (z-c_{K-2})^2)' \quad (9)$$

$$\varphi(\theta; \lambda, K) = \lambda \sum_{j=1}^P \left(\sum_{k=1}^K \theta_{j,k}^2 \right)^{1/2} \quad (10)$$

This is the model in the comparison whose sole job is to isolate the effect of univariate nonlinearity while holding additivity fixed: each characteristic is transformed nonlinearly but the model still sums across characteristics with no cross terms. Its comparative performance is therefore the paper's cleanest single piece of evidence on whether nonlinearity alone, absent interactions, is what the more flexible methods are exploiting.

3.6 Regression Trees, Random Forests, and Gradient Boosting

A regression tree partitions the predictor space into K rectangular regions and fits a constant in each,

$$g(z_{i,t}; \theta, K, L) = \sum_{k=1}^K \theta_k \cdot 1\{z_{i,t} \in C_k(L)\} \quad (11)$$

where the partition is built by recursive binary splitting to depth L , choosing at each node the single split that most reduces within-partition squared error. This is precisely the mechanism through which trees represent interactions: a leaf is defined by a conjunction of conditions on several characteristics simultaneously, for example, a region combining high momentum with small size, which no additive model can express without that exact interaction term being specified in advance. A random forest averages many such trees, each grown on a bootstrap resample with a randomly restricted subset of candidate splitting variables at each node ("dropout" of predictors), which decorrelates individual trees and reduces the variance of the average. Gradient boosting instead grows a sequence of shallow trees, each fit to the residuals left by the ensemble so far and added with a small shrinkage weight $\nu \in (0,1)$; depth L , shrinkage ν , and the number of trees B are the three parameters tuned on the validation sample.

3.7 Neural Networks

The feed-forward networks are built from the standard recursive layer definition,

$$x_k^{(l)} = \text{ReLU}(x^{(l-1)'} \theta^{(l-1)}_k), \quad \text{ReLU}(x) = \max(0, x) \quad (12)$$

$$g(z; \theta) = x^{(L-1)'} \theta^{(L-1)} \quad (13)$$

with five architectures compared, NN1 through NN5, of one to five hidden layers and, respectively, (32), (32,16), (32,16,8), (32,16,8,4), and (32,16,8,4,2) neurons per layer. Each successive architecture nests strictly more representational capacity, and because each hidden layer's output feeds nonlinearly into the next, a network with two or more layers can represent interaction effects among the raw predictors that a single linear or additive layer cannot. Overfitting risk is controlled through five simultaneous devices rather than one: an $L1$ penalty on the weights, early stopping against validation-sample loss, batch normalization between layers (Ioffe and Szegedy 2015), an adaptively decaying learning rate in the stochastic gradient optimizer (following Kingma and Ba's 2014 Adam-style scheme), and ensembling across ten networks trained from different random initializations. **CRITIQUE.** This layering of regularization devices is methodologically sound but makes it harder to attribute the networks' performance to any single mechanism; the paper's own architecture-depth results (Section 5.3) are the cleanest way to isolate capacity from the other four regularizers, since depth is varied while every other device is held fixed.

3.8 Out-of-Sample Evaluation

Performance is judged by an out-of-sample R-squared computed against a zero forecast rather than against the historical mean return,

$$R^2_{\text{oos}} = 1 - (\sum_{(i,t) \in T_3} (r_{i,t+1} - \hat{r}_{i,t+1})^2) / (\sum_{(i,t) \in T_3} r_{i,t+1}^2) \quad (14)$$

PAPER. The authors are explicit, and correct, that demeaning the denominator, as is conventional in some return-prediction studies, implicitly credits a model for out-of-sample predictability that it does not have, since the historical in-sample mean itself uses information a real-time forecaster would not yet possess in the same way; benchmarking against zero avoids that pitfall and is the more conservative, and more defensible, choice. This convention is one reason the reported R^2_{OOS} figures, though tiny in absolute percentage terms, should not be compared directly to demeaned in-sample R-squared values reported elsewhere in the literature.

04 Data and Empirical Strategy

4.1 Sample and Universe

The panel covers essentially the full universe of CRSP common stocks listed on the NYSE, AMEX, and NASDAQ from March 1957 to December 2016, roughly 30,000 distinct stocks over the sample, with an average of about 6,200 stocks per month. No minimum price or size filter is imposed beyond CRSP's own coverage rules, so microcaps and thinly traded names are included, a choice the authors defend as necessary to avoid pre-selecting toward the stocks for which linear models already work well, but one that the review returns to in Section 7 in connection with trading costs.

4.2 Predictors

The 94 characteristics are inherited from Green, Hand, and Zhang's (2017) audit of the published anomalies literature: 61 updated annually from Compustat fundamentals, 13 updated quarterly, and 20 updated monthly from price and volume data, together with 74 industry indicator variables built from the one-digit SIC code. Combined multiplicatively with the eight macro variables from Welch and Goyal (2008) and the constant, the total feature count is 920. **CRITIQUE.** Treating a curated, published list of 94 characteristics as the "raw" information set is a defensible practical choice but is not itself free of the multiple-testing problem the paper is otherwise careful about: each of the 94 characteristics survived its own individual publication process in a literature already shown to be prone to false discoveries (Harvey, Liu, and Zhu 2016; McLean and Pontiff 2016), so the paper's out-of-sample discipline protects against overfitting the machine learning stage but does not protect against having started from a predictor set that is itself the product of decades of specification search across the profession.

4.3 Sample Splitting

The identifying discipline of the entire exercise rests on a single design choice: eighteen years (1957 to 1974) are used purely for training, the following twelve years (1975 to 1986) purely for tuning hyperparameters against validation-sample loss, and the final thirty years (1987 to 2016) are held out purely for evaluation. The three windows are then rolled forward one year at a time, refitting annually, always keeping training strictly before validation strictly before the test observation being scored, so no information from a stock's own future ever enters its forecast. This is where the identifying variation in the paper's central claim comes from: because every method sees an identical, chronologically honest split, a difference in test-period R-squared

between two methods cannot be attributed to one method having quietly used more information than another, it can only reflect a difference in how efficiently each method extracts predictive structure from the same information.


4.4 From Stock Forecasts to Portfolios

Beyond scoring individual-stock forecasts, the paper converts predicted returns into two portfolio exercises: a bottom-up market-timing signal for the S&P 500 built by value-weighting individual stock forecasts, and decile portfolios formed monthly by sorting all stocks on their predicted return and holding the extreme-decile long-short spread, both value-weighted and equal-weighted. **INFERENCE.** The equal-weighted spread in particular gives very small stocks the same portfolio weight as large ones, which is informative about the theoretical strength of the identified signal but, without an explicit adjustment, is not by itself informative about the return achievable by a real investor operating at scale, a distinction the review develops further in Section 7.

05 Results

5.1 Out-of-Sample Predictive Accuracy

Table 1 of the paper reports monthly R^2_{OOS} at the individual-stock level. The unrestricted linear model is strongly negative, -3.46 percent overall and -11.28 percent among the largest 1,000 stocks, meaning that OLS with all 920 regressors actively destroys predictive accuracy relative to a naive zero forecast, evidence of severe in-sample overfitting once collinear, uninformative interactions are added without penalty. The restricted three-characteristic linear benchmark recovers to a small positive 0.16 percent. Every machine learning method beats that linear benchmark, but not by equal amounts: penalized linear and dimension-reduction methods (elastic net, PCR, PLS, generalized linear model) cluster around 0.11 to 0.27 percent, while trees, boosting, and neural networks cluster higher, 0.33 to 0.40 percent, peaking at 0.40 percent for the three-layer network (NN3).


Figure 1. Monthly out-of-sample R^2 by method, full stock universe

Source: Gu, Kelly, and Xiu (2020), Table 1. Unrestricted OLS omitted from the axis range (-3.46%) to preserve legibility among the remaining methods; its bar is shown truncated.

EVIDENCE. The ordering is not a matter of scale, every method is trained on the identical 920-feature vector, so the roughly $2.5\times$ gap between the best penalized-linear method (0.27 percent, PLS) and the best neural network (0.40 percent, NN3) is attributable to functional form alone. **INFERENCE.** Because the generalized linear model, which adds univariate nonlinearity without interactions, performs no better than the penalized linear methods (0.19 percent), while trees and networks, which add interactions, perform substantially better, the paper's most defensible mechanistic claim, that the gain is coming specifically from interaction effects rather than from nonlinearity in general, is well supported by this contrast rather than merely asserted.

5.2 Economic Magnitude

A monthly R^2_{oos} under one percent sounds negligible in isolation, but individual monthly stock returns are extremely noisy, so the paper's own benchmarking convention (Section 3.8) makes even these figures large relative to what an efficient-markets prior would predict. The portfolio-level translation makes the magnitude concrete. Timing the S&P 500 with the bottom-up neural network forecast raises the annualized out-of-sample Sharpe ratio to 0.77, from 0.51 for buy-and-hold, a 26-percentage-point improvement from the same underlying signal. The extreme-decile long-short spread sorted on NN3 forecasts earns an annualized Sharpe ratio of 1.35 value-weighted and 2.45 equal-weighted, against 0.61 and 0.83 respectively for the OLS-3 benchmark.


Figure 2. Long-short decile portfolio Sharpe ratio, OLS-3 versus NN3

Source: Gu, Kelly, and Xiu (2020), Section 3 (portfolio results). Figures are gross of transaction costs; see Section 7.

CRITIQUE. The equal-weighted spread's Sharpe ratio of 2.45 is the single largest number reported anywhere in the paper, and it is also the number most exposed to the microcap and turnover concerns raised in Section 4.4 and revisited in Section 7; a reader should weight the value-weighted figure of 1.35 more heavily when judging what is realistically capturable, and even that figure should be read as gross of the trading costs a real interaction-heavy, momentum-loaded strategy would incur.

5.3 Depth and the Shallow-Learning Result

Within the neural-network family, accuracy rises from NN1 (0.33 percent) through NN2 (0.39 percent) to a peak at NN3 (0.40 percent), then declines at NN4 (0.39 percent) and NN5 (0.36 percent).

Figure 3. Monthly out-of-sample R^2 by network depth, NN1–NN5

Source: Gu, Kelly, and Xiu (2020), Table 1, full stock universe.

PAPER. The authors attribute this hump shape to the low signal-to-noise ratio and modest effective sample size of financial panels relative to the domains, image and speech recognition, where very deep networks were developed, concluding that "shallow learning" outperforms "deep learning" in this setting. **INFERENCE.** This is a genuinely falsifiable, and here falsified, prediction: if depth were simply a free source of representational power, accuracy would rise monotonically with the number of layers, and it does not.

5.4 Variable Importance

Across methods, the characteristics carrying the most predictive weight fall into three families: price-trend measures (stock momentum, industry momentum, short-term reversal), liquidity measures (market value, dollar trading volume, bid-ask spread), and volatility measures (idiosyncratic and total return volatility, market beta, beta squared). **INFERENCE.** This ranking connects the paper directly, rather than merely by citation, to three of the empirical asset-pricing literature's most established results: momentum (Jegadeesh and Titman 1993), illiquidity pricing (Amihud 2002), and volatility-related anomalies, which is independent, method-agnostic corroboration that the machine learning stage is recovering economically recognizable structure rather than an artifact specific to any one estimator.

06 Robustness and Additional Tests

CONCERN: HEAVY-TAILED RETURNS DISTORT SQUARED-ERROR TRAINING

Replacing squared loss with Huber loss for elastic net, the generalized linear model, and gradient boosting. Result: each Huber variant modestly outperforms its squared-loss counterpart. This resolves the concern for the methods tested, though the paper does not report a Huber variant for every method in the comparison, so the robustness is demonstrated rather than exhaustive.

CONCERN: GAINS ARE A SMALL-STOCK, ILLIQUIDITY-DRIVEN ARTIFACT

Re-estimating separately on the largest 1,000 and smallest 1,000 stocks. Result: nonlinear methods' advantage over linear methods is, if anything, larger among large, liquid stocks (NN3 reaches 0.70 percent among the top 1,000 versus 0.45 percent among the bottom 1,000). This meaningfully resolves the narrow concern that the headline result is a large-stock-versus-small-stock artifact, since the advantage survives, and even strengthens, among the most liquid names. **CRITIQUE.** It does not, however, resolve the separate and more consequential concern that the equal-weighted portfolio Sharpe ratios reported in Section 5.2 are still dominated by small-stock exposure through equal weighting, which this size-split test does not directly address.

CONCERN: THE RESULT IS AN ARTIFACT OF THE ONE-MONTH FORECAST HORIZON

Re-running the comparison at an annual horizon. Result: annual R^2_{OOS} values are roughly an order of magnitude larger across all methods (for example, about 3.4 percent overall, versus 0.33 to 0.40 percent monthly, for the leading methods), with a broadly similar ranking across methods. This is a sensible mechanical consequence of averaging idiosyncratic noise over a longer horizon rather than independent confirmation of a different economic mechanism, and it partially, not fully, resolves the horizon-dependence concern, since it confirms the ranking is stable but does not by itself validate the annual-horizon economic magnitude to the same standard as the monthly result.

CONCERN: THE LINEAR-VERSUS-NONLINEAR RANKING COULD REVERSE UNDER A DIFFERENT, EQUALLY PLAUSIBLE DATA-GENERATING PROCESS

A Monte Carlo exercise (Appendix A) simulating returns under two designs, one linear and additive, one nonlinear with predictor interactions, and applying the full method menu to each simulated panel. Result: linear and generalized additive methods dominate under the linear design; trees and neural networks dominate under the interactive design. This directly resolves the concern, because it shows the ranking observed in the real data is exactly what the same methods produce when the truth is known to require interactions, and the reverse when it does not, which is the strongest single piece of evidence in the paper that functional form, not some artifact of the empirical panel, is driving the result.

07 Critical Evaluation

7.1 What the Paper Establishes

ESTABLISHED

Holding the information set fixed, methods capable of representing interactions among characteristics achieve materially higher out-of-sample predictive accuracy for individual stock returns than additive linear or nonlinear methods, and this ranking reproduces in a controlled simulation where the true data-generating process is known.

ESTABLISHED

The dominant predictive characteristics, momentum, liquidity, and volatility measures, are consistent across all nine method families, which rules out the concern that the ranking in 7.1 above is an artifact specific to one estimator's idiosyncrasies.

SUPPORTED BUT CONDITIONAL

The translation of stock-level forecast accuracy into large portfolio Sharpe ratios (1.35 value-weighted, 2.45 equal-weighted) holds under the paper's own construction, no trading costs, monthly rebalancing, decile extremes, and is conditional on that construction rather than on a realized, cost-adjusted, capacity-constrained implementation.

SUGGESTIVE

That the recovered structure reflects genuine, stable risk-return trade-offs rather than a decaying, previously documented set of anomalies interacting in a way linear models could not exploit is consistent with the evidence but not separately established; the paper does not test whether its forecasts retain power after excluding stocks and periods in which momentum, size, and volatility anomalies are already known to concentrate.

NOT ESTABLISHED

Why the interactions exist economically, that is, which equilibrium mechanism generates a return premium that depends jointly on, for instance, small size and high momentum rather than on either characteristic alone, is explicitly outside the paper's scope, and the authors say so directly.

7.2 Strengths

Recognized by the authors: the shared feature set and identical, chronologically honest evaluation protocol across nine method families; the Monte Carlo validation of the interaction mechanism; the cross-method consistency of variable importance as a check against overfitting.

Identified independently: the paper's choice to benchmark R^2_{oos} against zero rather than the historical mean (Section 3.8) is a small technical decision with an outsized effect on the credibility of the entire numerical exercise, and it is executed correctly and defended transparently. The depth experiment within the neural-network family (Section 5.3) is an unusually clean natural test, because depth is varied while every other design choice, features, regularization, optimizer, is held fixed, which is rarer in machine learning applications than it should be.

7.3 Weaknesses and Independent Critique

PROBLEM: TRANSACTION COSTS AND CAPACITY ARE NOT MODELED

Why it matters: the leading signals identified, short-term reversal, momentum, small size, are precisely the characteristics associated with high portfolio turnover and elevated trading costs in the broader literature; a strategy sorted on these signals monthly will trade far more than a static value or size portfolio. Potential bias: reported Sharpe ratios, especially the equal-weighted 2.45 figure, likely overstate what is achievable net of costs and at meaningful scale. How it could be tested: applying a standard proportional or size-dependent transaction-cost model to the decile turnover implied by the paper's own predicted-return rankings, an exercise the paper does not perform and that subsequent research (see Section 9) has taken up directly.

PROBLEM: THE 94-CHARACTERISTIC UNIVERSE IS ITSELF A PRODUCT OF SPECIFICATION SEARCH

Why it matters: Green, Hand, and Zhang's (2017) list, which this paper treats as its raw information set, was itself extracted from a much larger set of published anomalies, each individually exposed to the multiple-testing concerns documented by Harvey, Liu, and Zhu (2016). Potential bias: the paper's rigorous out-of-sample discipline protects the machine-learning stage from overfitting but cannot retroactively correct for selection bias that entered at the stage where the 94 characteristics were chosen by the profession. How it could be tested: re-running the full comparison on a characteristic set constructed independently of the published-anomaly literature, for instance from raw accounting and price data without characteristic-level curation, a test genuinely difficult to execute cleanly and one the paper does not attempt.

PROBLEM: NO DIRECT TEST OF ANOMALY DECAY

Why it matters: McLean and Pontiff (2016) show that published anomalies lose roughly half their return predictability after publication, as trading activity arbitrages them away; because the paper's sample runs through 2016 and its leading signals are well-known, published anomalies, the reported gains could in principle partly reflect a period before full arbitrage rather than a permanent feature of the cross-section. Potential bias: forward-looking extrapolation of the reported Sharpe ratios risks overstating future performance. How it could be tested: splitting the test period into pre- and post-publication-intensity subperiods for the dominant characteristics and checking whether the machine learning advantage over linear methods narrows, a test the paper's own rolling, annually refit design would support without much additional machinery, but which it does not report.

SINGLE MOST IMPORTANT WEAKNESS

Of these, the absence of any transaction-cost adjustment is the most consequential, because it is the one weakness that bears directly on the paper's own headline economic-significance claim rather than on a secondary robustness question, and because the signals responsible for the largest gains are exactly the ones for which trading costs are known, from independent evidence elsewhere in the literature, to be economically material.

08 Contradictions and Edge Cases

EXTENSION. Under what conditions should the nonlinear advantage disappear? The Monte Carlo evidence already answers this directly: under a linear, non-interactive data-generating process, the advantage not only disappears but reverses, additive methods dominate. The empirical question this leaves open is whether the true return-generating process is stationary in this respect. If the interactions the paper recovers are themselves a byproduct of a specific market regime, elevated volatility clustering around 2000 to 2002 and 2008 to 2009, for instance, then a data-generating process that is closer to additive in calmer regimes would erode the nonlinear methods' advantage precisely when it is measured over calm subperiods, a decomposition the paper does not report.

EXTENSION. Does the mechanism survive at extreme values of network depth? Section 5.3 already shows the answer is no beyond three layers, and the decline from NN3 to NN5 is itself evidence against an unbounded "more capacity is better" reading of the results; the paper's own preferred explanation, low signal-to-noise ratio relative to sample size, is consistent with this decline but is not independently verified against an alternative explanation, that the optimization and regularization scheme tuned for shallower networks is simply less well suited to the deeper ones.

EXTENSION. Could a different mechanism generate the same cross-method ranking? Because trees, boosting, and neural networks share the capacity to represent interactions but differ substantially in how they regularize and optimize, an alternative explanation for their common outperformance is that they are simply harder to overfit against the validation sample than the additive methods, rather than that interactions specifically are the active ingredient. The Monte Carlo design (Section 6) is the paper's answer to exactly this alternative, and it is a reasonably strong answer, but it examines only two stylized data-generating processes rather than a continuum of interaction intensities, so it establishes that interactions can produce the observed ranking, not that they are the only structure that could.

EXTENSION. Does the conclusion hold outside the original setting? The panel is US common stocks only; whether the same ranking of methods, and the same three dominant signal families, would recur in other developed equity markets, in fixed income, or in currency cross-sections is an open empirical question the paper does not address and that later work (Section 9) has only partially taken up.

09 Further Literature

Table 1. Selected literature connections

Paper	Shared idea	Relationship
Jegadeesh and Titman (1993)	Momentum as a return predictor	Predecessor; momentum is the single most important characteristic recovered by every method in Gu, Kelly, and Xiu, independently corroborating the earlier linear finding under a fully nonlinear estimation framework.
Amihud (2002)	Illiquidity pricing	Predecessor; liquidity measures rank among the top predictive signals across all nine methods, again a nonlinear, cross-method corroboration of an established linear result.
Fama and French (2008); Lewellen (2015)	Parsimonious linear characteristic models	Direct linear benchmark; OLS-3 in this paper is Lewellen's three-characteristic specification, used explicitly as the yardstick every machine learning method is measured against.
Green, Hand, and Zhang (2017)	Systematic audit of the characteristics literature	Immediate empirical foundation; supplies the 94-characteristic predictor set this paper treats as its raw information set without re-deriving it.
Harvey, Liu, and Zhu (2016); McLean and Pontiff (2016)	Multiple testing and anomaly decay in the cross-section	Contradictory pressure; both papers raise concerns, data mining and post-publication decay, that this paper's out-of-sample protocol mitigates only partially, as discussed in Section 7.3.

Paper	Shared idea	Relationship
Kelly, Pruitt, and Su (2019)	Return predictability from firm characteristics	Competing methodology; imposes a latent linear factor structure (instrumented principal components) on characteristics rather than a flexible nonparametric function, trading Gu, Kelly, and Xiu's functional flexibility for an explicit risk-factor interpretation.
Kozak, Nagel, and Santosh (2020)	Shrinkage applied to characteristic-based factors	Competing methodology; applies Bayesian shrinkage to a linear stochastic discount factor built from characteristics, addressing the same overfitting concern with penalization rather than nonlinear functional form.
Freyberger, Neuhierl, and Weber (2020)	Nonparametric characteristic-return relationships	Same journal issue, contrasting design; adopts an additive nonparametric model, closely analogous to this paper's generalized linear model benchmark, and finds a similar limitation, that nonlinearity without interactions captures less than fully flexible methods.
Moritz and Zimmermann (2016)	Tree-based conditional portfolio sorts	Direct predecessor; the earliest systematic use of regression trees for exactly the interaction-capturing purpose this paper generalizes across a much wider method menu.
Chen, Pelger, and Zhu (2024, working paper 2019)	Deep learning for the stochastic discount factor	Direct extension; imposes a no-arbitrage constraint on a neural-network estimator of the pricing kernel, addressing this paper's own stated limitation that a forecasting improvement is not a structural asset-pricing model.

Compiled by the reviewer from the cited papers; not reproduced from Gu, Kelly, and Xiu (2020).

10 Research Gaps

PAPER-LEVEL GAPS

The paper does not report cost-adjusted portfolio performance, does not decompose predictive accuracy by market regime or volatility state, and does not test whether its forecasts retain power once known, previously documented anomaly signals are stripped from the feature set.

LITERATURE-LEVEL GAPS

The broader literature this paper sits within has not fully resolved whether machine-learning-based return predictability represents a genuinely new, previously inaccessible risk-return relationship or a more efficient extraction of the same risk premia the linear literature had already partially identified, nor whether the interaction structure recovered here is stable across international markets and asset classes outside US equities.

METHODOLOGICAL GAPS

Better data would mean characteristic panels constructed independently of the published-anomaly search process described in Section 7.3; better identification would mean a design that separates genuine interaction effects from the possibility that flexible methods are simply harder

to overfit against a validation sample; better measurement would mean reporting predictive accuracy jointly with realistic, characteristic-dependent transaction-cost and capacity estimates rather than as a separate exercise left to later papers.

11 Research Opportunities

1. Cost- and capacity-adjusted machine learning portfolios

RESEARCH QUESTION How much of the Sharpe-ratio advantage of nonlinear over linear return forecasts survives realistic, characteristic-dependent transaction costs and a constraint on assets under management?

MOTIVATION Section 7.3 identifies this as the paper's single most consequential unaddressed weakness.

HYPOTHESIS The nonlinear advantage narrows substantially, but does not vanish, once trading costs are applied, because part of the advantage comes from volatility and size characteristics that are less turnover-intensive than the momentum component.

ECONOMIC MECHANISM Portfolio turnover scales with the persistence of the underlying signal; short-horizon reversal-heavy signals imply high turnover and cost drag, while slower-moving volatility and size signals imply less.

LITERATURE CONNECTION Directly extends Gu, Kelly, and Xiu (2020); draws on the transaction-cost literature applied to characteristic-based anomalies.

DATA CRSP and the same 94-characteristic panel, augmented with bid-ask spread and volume-based cost proxies.

VARIABLES Predicted return deciles, realized turnover per decile, proportional and size-dependent cost estimates.

METHOD Replicate the paper's forecasting step, then apply a turnover-aware backtest with cost models varying by market-capitalization tier.

IDENTIFICATION STRATEGY Held-out, chronologically ordered test period as in the original paper, so costs are applied to genuinely out-of-sample decisions.

EXPECTED CONTRIBUTION A realistic, implementable estimate of the strategy's net-of-cost economic value.

MAIN RISK Reliable, characteristic-level cost data for the full 1987 to 2016 period, especially for microcap names, may not be available at the granularity required.

2. Stripping known anomalies to test for genuinely new structure

RESEARCH QUESTION Does the nonlinear methods' advantage over linear methods persist after orthogonalizing the feature set against the momentum, size, and volatility characteristics already known to drive the ranking?

MOTIVATION Section 7.1 flags this as a "suggestive but not established" claim in the original paper.

HYPOTHESIS Residual predictive power after orthogonalization is small but still favors nonlinear methods, indicating a genuine, if modest, interaction structure beyond the three dominant characteristic families.

ECONOMIC MECHANISM If interactions among second-tier characteristics (profitability, investment, leverage) carry independent information once the dominant signals are removed, that would point to a richer joint pricing structure than the top-line variable-importance results alone suggest.

LITERATURE CONNECTION Builds on Green, Hand, and Zhang (2017) and directly on Gu, Kelly, and Xiu (2020).

DATA The same 94-characteristic CRSP and Compustat panel.

VARIABLES Residualized characteristics after regressing out momentum, size, liquidity, and volatility measures.

METHOD Re-run the full nine-method comparison on the residualized feature set and compare the resulting R^2_{oos} ranking to the original.

IDENTIFICATION STRATEGY Same rolling, chronologically ordered train, validate, test split as the original paper.

EXPECTED CONTRIBUTION A direct test of whether the paper's result is fundamentally about three known anomalies interacting nonlinearly, or about a broader class of interaction effects.

MAIN RISK Orthogonalization can mechanically remove genuine information if the "known" characteristics are themselves noisy proxies for the true underlying signal.

3. Regime-conditional decomposition of nonlinear predictability

RESEARCH QUESTION Is the nonlinear methods' advantage over linear methods concentrated in high-volatility, recessionary periods, or is it stable across the business cycle?

MOTIVATION Raised directly in Section 8; the paper reports only full-sample averages.

HYPOTHESIS The advantage is larger in downturns, when characteristic interactions, for example the joint behavior of size and momentum during liquidity crunches, are economically more pronounced.

ECONOMIC MECHANISM Funding constraints and flight-to-quality effects during downturns can make characteristic interactions economically active in a way they are not during calm periods, consistent with intermediary-based asset-pricing mechanisms.

LITERATURE CONNECTION Connects Gu, Kelly, and Xiu (2020) to the intermediary asset-pricing literature and to NBER business-cycle dating.

DATA The same panel, split by NBER recession dates and by realized market volatility terciles.

VARIABLES R^2_{OOS} computed separately within each subperiod for each method.

METHOD Subsample re-estimation of the original comparison using the already-fitted, rolling annual models.

IDENTIFICATION STRATEGY Purely a within-test-period decomposition; no new estimation risk beyond the original paper's.

EXPECTED CONTRIBUTION Clarifies whether the paper's average result is a stable structural feature or a period-specific phenomenon concentrated in a few crisis episodes.

MAIN RISK A small number of recession months in a sixty-year sample limits statistical power for a regime split.

4. Cross-market replication outside US equities

RESEARCH QUESTION Does the same ranking of methods, and the same dominant signal families, momentum, liquidity, volatility, recur in non-US developed equity markets with comparable data availability?

MOTIVATION Section 8 identifies external validity as untested; this is prioritized over a generic "test it elsewhere" idea because the paper's own three dominant signals have well-documented, market-specific strength and weakness patterns internationally (momentum, in particular, is known to behave differently in some non-US markets), so the setting plausibly changes the economics rather than merely relabeling the exercise.

HYPOTHESIS The nonlinear advantage replicates in direction but is smaller in magnitude in markets with lower liquidity-characteristic dispersion.

ECONOMIC MECHANISM If interactions between liquidity and momentum drive much of the US result, markets with a narrower cross-section of liquidity should show a correspondingly narrower nonlinear advantage.

LITERATURE CONNECTION Connects to Rapach, Strauss, and Zhou (2013) on international return predictability and to the international asset-pricing literature more broadly.

DATA A comparable characteristic panel for, for example, the UK, Japan, and continental European markets, built from an equivalent to CRSP and Compustat coverage.

VARIABLES The same characteristic and macro set, adapted to available local data.

METHOD Identical nine-method comparison and rolling evaluation protocol.

IDENTIFICATION STRATEGY Same chronological train, validate, test discipline, market by market.

EXPECTED CONTRIBUTION Evidence on whether the interaction mechanism is a US-specific feature of market structure or a general property of equity cross-sections.

MAIN RISK Non-US characteristic and accounting data of comparable depth and length are considerably harder to assemble, and coverage gaps could be mistaken for a genuine economic difference.

5. A no-arbitrage-constrained extension of the comparison

RESEARCH QUESTION How much of the nonlinear forecasting advantage survives, and how does its economic interpretation change, when the estimated function is required to satisfy a no-arbitrage pricing-kernel restriction rather than left as an unconstrained forecast?

MOTIVATION Directly addresses the paper's own stated limitation (Section 3.8's companion discussion in the original text) that a forecasting improvement is a measurement, not a structural model.

HYPOTHESIS Imposing a no-arbitrage constraint reduces raw predictive accuracy modestly but yields a stochastic discount factor whose implied risk prices are more stable across subperiods than the unconstrained forecast.

ECONOMIC MECHANISM A constrained pricing kernel forces the interaction structure identified in the unconstrained forecast to be consistent with the absence of arbitrage across the full cross-section simultaneously, not stock by stock.

LITERATURE CONNECTION Directly continues Chen, Pelger, and Zhu (2024) and connects back to Kelly, Pruitt, and Su (2019).

DATA The same CRSP panel and characteristic set as the original paper, for comparability.

VARIABLES Estimated stochastic discount factor loadings in place of raw return forecasts.

METHOD A generative adversarial or constrained-optimization neural-network estimator of the pricing kernel, benchmarked against the original paper's unconstrained forecasts on the identical test period.

IDENTIFICATION STRATEGY Comparison on the same held-out 1987 to 2016 window, holding the feature set fixed to isolate the effect of the no-arbitrage constraint alone.

EXPECTED CONTRIBUTION A bridge between the paper's purely predictive framework and a genuinely structural asset-pricing interpretation of the same interaction structure.

MAIN RISK No-arbitrage-constrained neural estimators are computationally demanding and their optimization landscape is less well understood than the unconstrained case, raising replication risk.

Ranking

Idea 1 (cost- and capacity-adjusted portfolios) ranks first: it is the most feasible with existing public data, has the highest immediate economic importance because it bears directly on whether the paper's headline result is investable, and requires no new theoretical machinery. Idea 2 (stripping known anomalies) ranks second: it is methodologically straightforward, uses data already assembled by the original authors, and would materially sharpen the paper's own "suggestive but not established" finding into either an established or a rejected one. Idea 5 (no-arbitrage-constrained extension) ranks third: it is the most novel and the most consequential if successful, since it would convert the paper's purely predictive contribution into a structural one, but it carries the highest technical and computational risk of the five, which is why it is ranked below the two more immediately executable ideas.

12 Learning Framework

PREREQUISITE CONCEPTS

Ordinary least squares and ridge or lasso regularization; bias-variance trade-off; cross-validation and honest sample splitting; the basic mechanics of a regression tree and of gradient boosting; the feed-forward neural network and backpropagation; the concept of out-of-sample R-squared as distinct from in-sample fit.

CORE CONCEPTS LEARNED

That functional form, not the information set, can be the binding constraint on measured predictability; that interactions among predictors, not nonlinearity in general, are the specific feature that flexible methods exploit here; that model capacity has a demonstrable optimum, not a monotonic benefit, in low signal-to-noise financial panels; and that a rigorous, honest evaluation protocol is itself a research contribution independent of any single method's performance.

IMPORTANT EQUATIONS

The additive decomposition of returns into a conditional expectation and an unforecastable residual (equations 1 to 2); the elastic net penalty nesting lasso and ridge (equation 6); the tree partition representation (equation 11), which is the cleanest formal statement of how a nonlinear method encodes an interaction; and the zero-benchmarked out-of-sample R-squared (equation 14), whose construction materially affects how the paper's headline numbers should be read.

METHODS LEARNED

Elastic net regularization; principal components and partial least squares as dimension-reduction regressions; group lasso over spline bases as a route to additive nonlinearity; regression tree, random forest, and gradient boosted tree construction; feed-forward neural network architecture design and its regularization toolkit (L1 penalization, early stopping, batch normalization, adaptive learning rates, ensembling); Monte Carlo simulation as a device for validating a proposed mechanism against alternative data-generating processes.

QUESTIONS TO TEST UNDERSTANDING

Comprehension. In your own words, why does an unrestricted linear regression on 920 features produce a negative out-of-sample R-squared while a three-variable linear regression produces a positive one?

Analysis. Why does the generalized linear model with spline-based univariate nonlinearity fail to close the gap with tree-based and neural-network methods, and what does that failure specifically rule in or rule out about the source of the paper's central result?

Research. Design a single additional test, beyond the five robustness checks in Section 6, that would most sharply distinguish "the gain is from interactions among a stable set of characteristics" from "the gain is from a flexible method's superior ability to avoid overfitting a fixed set of characteristics, whether or not they truly interact."

PAPERS TO READ NEXT

Freyberger, Neuhierl, and Weber (2020), for the cleanest single contrast between additive nonparametric and fully interactive methods; Kelly, Pruitt, and Su (2019), for the competing latent-factor route to the same cross-sectional prediction problem; and Chen, Pelger, and Zhu (2024), for the natural next step of imposing a no-arbitrage structure on a neural-network pricing model.

13 Conclusion

ONE-SENTENCE THESIS

Holding the information available to a forecaster fixed, the choice of functional form, specifically, whether the estimator can represent interactions among characteristics, is the dominant determinant of how much of the cross-section of individual stock returns is measurably predictable.

THREE MOST IMPORTANT INSIGHTS

First, unrestricted linear regression is not a safe default when the feature space is large; it can perform worse than a naive forecast. Second, the specific mechanism separating strong from weak performers among the flexible methods is the capacity to model interactions, not flexibility or nonlinearity in the abstract, a distinction the paper isolates unusually cleanly through the generalized linear model contrast and the Monte Carlo exercise. Third, model capacity in this domain has an interior optimum, three hidden layers here, rather than a monotonic benefit, a finding with direct implications for how future work in this area should be designed.

STRONGEST EVIDENCE

The Monte Carlo simulation (Section 6), because it is the one piece of evidence in the paper that does not depend on any feature of the real, historically contingent stock-return panel, and it reproduces the empirical ranking exactly when interactions are known, by construction, to be present or absent.

BIGGEST WEAKNESS

The absence of a transaction-cost- and capacity-adjusted performance measure, which leaves the paper's economic-significance claim, its most attention-grabbing result, resting on gross rather than net portfolio returns (Section 7.3).

MOST IMPORTANT RESEARCH GAP

Whether the recovered predictability is a stable structural feature of the cross-section or a still-decaying remnant of already-published anomalies interacting in a way linear models could not previously exploit (Section 10).

BEST RESEARCH OPPORTUNITY

A cost- and capacity-adjusted re-evaluation of the paper's own portfolio construction, because it is immediately executable with the same public data, addresses the single most consequential weakness identified in this review, and would settle, rather than merely qualify, the paper's central economic claim.

THREE PAPERS TO READ NEXT

Freyberger, Neuhierl, and Weber (2020); Kelly, Pruitt, and Su (2019); Chen, Pelger, and Zhu (2024).

14 References

- Amihud, Yakov. 2002. Illiquidity and Stock Returns: Cross-Section and Time-Series Effects. *Journal of Financial Markets* 5(1): 31–56.
- Chen, Luyang, Markus Pelger, and Jason Zhu. 2024. Deep Learning in Asset Pricing. *Management Science* 70(2): 714–750.
- Fama, Eugene F., and Kenneth R. French. 2008. Dissecting Anomalies. *Journal of Finance* 63(4): 1653–1678.
- Freyberger, Joachim, Andreas Neuhierl, and Michael Weber. 2020. Dissecting Characteristics Nonparametrically. *Review of Financial Studies* 33(5): 2326–2377.
- Green, Jeremiah, John R. M. Hand, and X. Frank Zhang. 2017. The Characteristics that Provide Independent Information about Average U.S. Monthly Stock Returns. *Review of Financial Studies* 30(12): 4389–4436.
- Gu, Shihao, Bryan T. Kelly, and Dacheng Xiu. 2020. Empirical Asset Pricing via Machine Learning. *Review of Financial Studies* 33(5): 2223–2273.
- Harvey, Campbell R., Yan Liu, and Heqing Zhu. 2016. ... and the Cross-Section of Expected Returns. *Review of Financial Studies* 29(1): 5–68.
- Hutchinson, James M., Andrew W. Lo, and Tomaso Poggio. 1994. A Nonparametric Approach to Pricing and Hedging Derivative Securities via Learning Networks. *Journal of Finance* 49(3): 851–889.
- Ioffe, Sergey, and Christian Szegedy. 2015. Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift. *Proceedings of the 32nd International Conference on Machine Learning*.
- Jegadeesh, Narasimhan, and Sheridan Titman. 1993. Returns to Buying Winners and Selling Losers: Implications for Stock Market Efficiency. *Journal of Finance* 48(1): 65–91.
- Kelly, Bryan T., Seth Pruitt, and Yinan Su. 2019. Characteristics Are Covariances: A Unified Model of Risk and Return. *Journal of Financial Economics* 134(3): 501–524.
- Kingma, Diederik P., and Jimmy Ba. 2014. Adam: A Method for Stochastic Optimization. *arXiv working paper 1412.6980*.
- Kozak, Serhiy, Stefan Nagel, and Shrihari Santosh. 2020. Shrinking the Cross-Section. *Journal of Financial Economics* 135(2): 271–292.

- Lewellen, Jonathan. 2015. The Cross-section of Expected Stock Returns. *Critical Finance Review* 4(1): 1–44.
- McLean, R. David, and Jeffrey Pontiff. 2016. Does Academic Research Destroy Stock Return Predictability? *Journal of Finance* 71(1): 5–32.
- Moritz, Benjamin, and Tom Zimmermann. 2016. Tree-Based Conditional Portfolio Sorts: The Relation between Past and Future Stock Returns. Working paper, Ludwig Maximilian University of Munich and University of Cologne.
- Rapach, David E., Jack K. Strauss, and Guofu Zhou. 2013. International Stock Return Predictability: What Is the Role of the United States? *Journal of Finance* 68(4): 1633–1662.
- Welch, Ivo, and Amit Goyal. 2008. A Comprehensive Look at the Empirical Performance of Equity Premium Prediction. *Review of Financial Studies* 21(4): 1455–1508.