

Noise Dressing of Financial Correlation Matrices

Laurent Laloux, Pierre Cizeau, Jean-Philippe Bouchaud, Marc Potters

PUBLICATION YEAR	1999 (arXiv preprint submitted October 1998)
JOURNAL	Physical Review Letters, volume 83, issue 7, page 1467
DOI	10.1103/PhysRevLett.83.1467
RESEARCH AREA	Random matrix theory applied to portfolio risk
SERIES	Portfolio Management
PAPER TYPE	Empirical and methodological (mixed)
REVIEWED BY	Arthur Faugeron
DATE REVIEWED	29 September 2026

ANALYTICAL ABSTRACT

The letter asks whether the empirical correlation matrix of 406 S&P 500 stocks, estimated from 1309 daily observations between 1991 and 1996, contains anything beyond sampling noise. Its benchmark is the eigenvalue density of a random Wishart matrix built from independent series, which random matrix theory fixes in closed form and which has sharp edges at $\sigma^2(1 + 1/Q \pm 2\sqrt{1/Q})$, where $Q = T/N$. Fitting this density to the bulk of the empirical spectrum leaves only about 6 percent of the eigenvalues above the upper edge, led by a market mode roughly 25 to 30 times larger than that edge. Eigenvector components inside the noise band are statistically indistinguishable from a Gaussian, whereas those above it are not. The authors conclude that the small eigenvalues, which govern minimum-risk Markowitz portfolios, are dominated by measurement noise, which casts doubt on the blind use of raw empirical correlation matrices for risk management, and they propose the same benchmark as a tool to isolate genuine correlations. This review reconstructs the mathematics, reconciles the numbers reported in the letter, and tests its claims with simulations.

1. Introduction

Source note. The paper is a three-page letter. This review was built from the arXiv full text (cond-mat/9810255, opened and read in this session) and its listing metadata, which gives the journal reference and DOI. The typeset Physical Review Letters version was not consulted, and the PDF could not be downloaded into the library because the retrieval environment blocked direct file transfer. Every quantitative statement attributed to the authors below comes from that text. Where this review computes a number itself, the text says so and labels it as a calculation or a simulation. The letter contains two figures and no tables; every table and every other figure in this review is the reviewer's own construction.

1.1 What is the paper trying to discover?

The paper asks how much of the structure in a sample correlation matrix of stock returns is information about the economy, and how much is the arithmetic residue of estimating $N(N-1)/2$ correlations from finite series. Its operational question is sharper: does the eigenvalue spectrum of the empirical correlation matrix of 406 stocks differ from the spectrum that a matrix of pure noise of the same dimensions would produce, and if so, in which part of the spectrum?

1.2 Why does the question matter, and what was known?

Markowitz (1952) portfolio selection requires the covariance matrix as its central input, and the minimum-variance solution loads heavily on the directions in which returns co-vary least, that is, on the smallest eigenvalues. Practitioners had long known that sample covariance matrices are unreliable when the number of assets is comparable to the number of observations (Elton and Gruber, 1973; Jobson and Korkie, 1980; Michaud, 1989), and the standard remedy was to impose structure through factor or constant-correlation models. What was missing was a quantitative, parameter-light description of what pure noise looks like in the spectrum, so that noise could be told apart from signal directly.

1.3 What does the paper propose, and how?

The authors import a result from mathematical physics: the Marchenko-Pastur law (Marchenko and Pastur, 1967), which gives the limiting eigenvalue density of the sample covariance matrix of independent variables when N and T grow with the ratio $Q = T/N$ fixed. Their null hypothesis is that all assets are uncorrelated. They compare the predicted density and the predicted distribution of eigenvector components with those of an actual S&P 500 sample, and read departures from the null as information.

1.4 What does it find, and why does it matter?

About 94 percent of the empirical eigenvalues lie where the null predicts, while the remaining 6 percent lie above the upper edge, the largest by an order of magnitude. Eigenvectors inside the band look random, those outside do not. The finding matters because it converts a folk worry (noisy covariance matrices) into a measurable statement: in this sample most of the spectrum, including the whole low-risk end, is statistically noise. This underlies a large later literature on covariance cleaning.

2. Research Problem and Hypothesis

2.1 Research question

The paper asks whether the eigenvalue density and the eigenvector component distribution of the sample correlation matrix of 406 US stocks over 1309 trading days are consistent with those of a random matrix built from independent, identically distributed entries, and, where they are not, which eigenvalues carry the excess.

2.2 Hypotheses

The null hypothesis is that the normalised price changes $M_{it} = \delta x_i(t)$ are independent and identically distributed with unit variance, so that every off-diagonal true correlation is zero and all observed structure is sampling error. The alternative is that some eigenvalues, and their eigenvectors, depart from the null prediction, indicating genuine collective correlations. A refined null is also considered: the market mode is genuine, and the orthogonal complement of the market direction is pure noise with reduced variance $\sigma^2 = 1 - \lambda_1/N$.

2.3 Motivation and research gap

Portfolio construction needs a correlation matrix, the number of estimated parameters grows as N^2 , and the number of usable observations grows only as $N \cdot T$. In the sample used, $N(N-1)/2 = 82,215$ pairwise correlations are estimated from $N \cdot T = 531,454$ return observations, roughly 6.5 observations per parameter (calculation from the paper's N and T). The gap is the absence of an explicit, closed-form description of the noise floor that lets a researcher separate the two components without imposing a structural model.

2.4 Contribution and its evaluation

PAPER The paper contributes a benchmark and a diagnostic, not a new estimator. **EVIDENCE** It demonstrates for one sample that the benchmark fits most of the spectrum and that eigenvectors inside the band are Gaussian-like. **CRITIQUE** It does not demonstrate, by any portfolio calculation, that using the raw matrix worsens risk control, nor does it test the null statistically. The contribution is therefore real but narrower than the concluding language on risk management suggests.

3. Theoretical Framework and Mechanism

The framework is a null model from random matrix theory. The mathematics is reproduced below with the equation numbers of this review; the letter numbers four equations, which correspond to Equations (1), (2), (3) with (4), and (8) here. Equations (5), (6), (7), (9) and (10) are the reviewer's derivations, added to make the argument followable.

3.1 The empirical object

Let M be the $N \times T$ matrix of demeaned price changes, each series rescaled to constant unit volatility. The empirical correlation matrix and its eigenvalue density are

$$C = (1/T) M M^T, \quad C_{ij} = (1/T) \sum_{t=1}^T \delta x_i(t) \delta x_j(t) \quad (1)$$

$$\rho_C(\lambda) = (1/N) dn(\lambda)/d\lambda, \quad n(\lambda) = \text{number of eigenvalues of } C \text{ below } \lambda \quad (2)$$

N	number of assets (406)
T	number of daily observations (1309)
Q	T/N , observations per asset (3.22)
M_{it}	normalised price change of asset i on day t
λ_α, v_α	eigenvalue and unit eigenvector of C , $\alpha = 1, \dots, N$

Equation (1) states that C is a sample second-moment matrix. Because each series has unit variance, the diagonal is one and $\text{trace}(C) = N$, so the eigenvalues average to exactly one. Equation (2) is the empirical spectral distribution, which is the object compared with theory. When N is close to T , C is far from its population counterpart even when the population is the identity, and Equations (3) and (4) quantify exactly how far.

3.2 The null prediction: the Marchenko-Pastur law

If the entries of M are independent with mean zero and variance σ^2 , then for $N, T \rightarrow \infty$ with $Q = T/N \geq 1$ fixed the eigenvalue density of C converges to

$$\rho_C(\lambda) = [Q/(2\pi\sigma^2)] \cdot \sqrt{[(\lambda_{\max} - \lambda)(\lambda - \lambda_{\min})]}/\lambda, \quad \lambda_{\min} \leq \lambda \leq \lambda_{\max} \quad (3)$$

$$\lambda_{\max/\min} = \sigma^2 (1 + 1/Q \pm 2\sqrt{1/Q}) = \sigma^2 (1 \pm \sqrt{1/Q})^2 \quad (4)$$

The square-root numerator produces two sharp edges, the $1/\lambda$ factor concentrates density near the lower edge, and the density is zero outside $[\lambda_{\min}, \lambda_{\max}]$. The key economic implication is that pure noise generates a whole band of eigenvalues, from well below one to well above one, even though every true eigenvalue equals one. A researcher who reads the largest sample eigenvalues as strong factors, or the smallest as nearly riskless combinations, is reading the tails of this band.

Derivation sketch. With $c = 1/Q = N/T$, the Stieltjes transform $m(z) = \int \rho_c(\lambda)/(\lambda - z) d\lambda$ of the limiting density satisfies the quadratic

$$c \sigma^2 z m(z)^2 + (z - \sigma^2(1 - c)) m(z) + 1 = 0 \quad (5)$$

whose solution is $m(z) = [-(z - \sigma^2(1 - c)) + \sqrt{(z - \sigma^2(1 - c))^2 - 4c\sigma^2 z}] / (2c\sigma^2 z)$. The density is recovered as $\rho_c(\lambda) = \text{Im } m(\lambda + i0)/\pi$, which is non-zero only where the discriminant is negative. The discriminant $(\lambda - \sigma^2(1 - c))^2 - 4c\sigma^2 \lambda$ factorises as $(\lambda - \lambda_{\min})(\lambda - \lambda_{\max})$, with the roots in Equation (4), and substituting back gives Equation (3) exactly, since $1/c = Q$. Two checks are useful. The mean of the density is σ^2 , so the trace identity holds. As $Q \rightarrow 1$, $\lambda_{\min} \rightarrow 0$ and the density diverges as $1/\sqrt{\lambda}$, which the letter notes; as $Q \rightarrow \infty$, both edges collapse to σ^2 and C converges to the population matrix.

Numerical edges (calculation). For $Q = 1309/406 = 3.2241$, Equation (4) gives $\lambda_{\max} = 2.424$ and $\lambda_{\min} = 0.196$ for $\sigma^2 = 1$, and 2.060 and 0.167 for $\sigma^2 = 0.85$, and 1.794 and 0.145 for $\sigma^2 = 0.74$. A simulated pure-noise matrix of exactly this size (seed fixed, this review) has extreme eigenvalues of 0.200 and 2.395, inside the theoretical band, as expected at this N .

3.3 Refined null: a genuine market factor

The market mode cannot be noise. The authors therefore remove it and posit that the rest of the spectrum is noise with a smaller variance. Since the eigenvalues of a correlation matrix sum to N , the variance available to the noise bulk is what remains after the largest eigenvalue:

$$\sigma^2 = 1 - \lambda_1/N \quad (6)$$

The letter reports $\sigma^2 = 0.85$, which implies $\lambda_1 \approx 0.15 N \approx 61$ (calculation). This reconciles a point that reads as a discrepancy in the letter. The text says the top eigenvalue is 25 times larger than the predicted $\lambda_{\max}$ and the figure caption says 30 times. With $\lambda_1 \approx 61$, the ratio is $61/2.424 \approx 25$ against the $\sigma^2 = 1$ edge and $61/2.060 \approx 30$ against the $\sigma^2 = 0.85$ edge. **INFERENCE** The two statements are therefore consistent and differ only in the reference edge, which the letter does not make explicit.

The best fit is obtained for $\sigma^2 = 0.74$, lower than 0.85. Because the mean of the Marchenko-Pastur density is σ^2 , a bulk with $\sigma^2 = 0.74$ that holds 94 percent of the eigenvalues accounts for about 74 percent of total variance. This is exactly the statement in the letter that the top 6 percent of eigenvectors carry 26 percent of the total volatility ($1 - 0.74 = 0.26$). The market mode holds about 15 percent, so the remaining eigenvalues above the edge, roughly 23 of the 24 in total, share about 11 percent of variance, an average of roughly 1.9 (range 1.85 to 2.0 given the rounding of 0.85), only slightly above the fitted edge of 1.79. **INFERENCE** This suggests that the non-market information is present but modest in size, and the letter's description of these eigenvalues as exceeding the edge by a substantial amount should be read for the market mode rather than for the whole 6 percent.

3.4 Eigenvector test: the Porter-Thomas distribution

If the eigenvectors of a noise matrix carry no information, they are uniformly distributed on the unit sphere. Normalising components so that $\sum_i u_i^2 = N$ and imposing only the constraint that u has unit second moment, the maximum-entropy density solves

$$\max_P -\int P(u) \ln P(u) du \quad \text{subject to} \quad \int P(u) du = 1, \quad \int u^2 P(u) du = 1 \quad (7)$$

The Lagrangian gives $\ln P(u) = -1 - a - b u^2$, hence $P(u) \propto \exp(-b u^2)$, and the unit-variance constraint fixes $b = 1/2$:

$$P(u) = (2\pi)^{-1/2} \exp(-u^2/2) \quad (8)$$

Equation (8) is the null distribution for eigenvector components. The test is graphical: histogram the components of eigenvectors with eigenvalue inside $[\lambda_{\min}, \lambda_{\max}]$ and compare with Equation (8). An eigenvector that carries a factor, such as the market, has components concentrated near a common positive value and does not match. Because Equation (8) is also the large- N law for the components of uniformly random directions, agreement says that these directions cannot be distinguished from arbitrary rotations, not that they are proved to contain no information (see Section 7).

3.5 Why the low end of the spectrum matters: the Markowitz link

The letter states, without formulas, that the smallest eigenvalues correspond to the least risky portfolios and are dominated by noise. The mechanism can be written explicitly. The global minimum-variance portfolio for a covariance matrix C solves $\min_w w^T C w$ subject to $w^T 1 = 1$, with solution and risk

$$w^{opt} = C^{-1} 1 / (1^T C^{-1} 1), \quad R^2 = 1 / (1^T C^{-1} 1) = [\sum_{\alpha} (v_{\alpha} \cdot 1)^2 / \lambda_{\alpha}]^{-1} \quad (9)$$

Each eigen-direction enters the denominator with weight $1/\lambda_{\alpha}$, so the smallest sample eigenvalues dominate the optimiser. If those eigenvalues are pulled down by sampling error, the optimiser sees riskless combinations that do not exist. **EXTENSION** For a true covariance equal to the identity and $Q > 1$, known results (see Pafka and Kondor, 2004, and the simulation in Figure 3) give the in-sample and out-of-sample variances of the estimated minimum-variance portfolio relative to the true optimum $1/N$:

$$R_{in}^2 = (1/N)(1 - 1/Q), \quad R_{out}^2 = (1/N) \cdot 1/(1 - 1/Q), \quad R_{out}^2 / R_{in}^2 = 1/(1 - 1/Q)^2 \quad (10)$$

At $Q = 3.22$, Equation (10) gives an in-sample variance 31 percent below the true optimum, an out-of-sample variance 45 percent above it, and an out-of-sample to in-sample variance ratio of 2.10, so that realised volatility would exceed predicted volatility by a factor of about 1.45. As $Q \rightarrow 1$ the in-sample risk goes to zero and the out-of-sample risk diverges, which is the singularity of the density at $\lambda_{\min} = 0$. This calculation is not in the letter; it is the quantitative content of its verbal warning, valid only under the stated identity-covariance benchmark.

4. Data and Empirical Strategy

4.1 Research design and identifying variation

The design is a comparison of one realised spectrum with a theoretical spectrum, not a regression. The identifying variation is the gap between the empirical eigenvalue density and the Marchenko-Pastur curve, together with the gap between empirical eigenvector components and Equation (8). No causal claim is made or possible: the exercise establishes which parts of the sample correlation structure are statistically distinguishable from noise, not what economic forces produce the distinguishable parts. The single fitted parameter is σ^2 , adjusted to match the bulk of the histogram, so the evidence has one degree of freedom of tuning and a visual, rather than test-based, criterion of fit.

4.2 Data

Table 1. Data and parameters reported in the letter, with the reviewer's derived quantities

Item	Value	Status
Universe	S&P 500 stocks, N = 406	Paper
Sample period	1991 to 1996	Paper
Frequency	Daily price changes	Paper
Observations per series	T = 1309 trading days	Paper
Ratio Q = T/N	3.22	Paper
Preprocessing	Mean subtracted; each series rescaled to unit constant volatility	Paper
Pairwise correlations estimated	$N(N-1)/2 = 82,215$	Calculation
Return observations	$N \cdot T = 531,454$	Calculation
Observations per estimated correlation	about 6.5	Calculation
Selection rule for the 406 stocks	Not stated in the available text	Not established
Price source and adjustments	Not stated in the available text	Not established

Source: Laloux, Cizeau, Bouchaud and Potters (1999) for rows marked Paper; reviewer's arithmetic otherwise.

The unit of observation is a stock-day. Normalising each series to constant unit volatility removes cross-sectional differences in volatility, so the matrix analysed is a correlation matrix, and it is also stated that the same analysis was repeated on a matrix of volatility correlations built from intraday fluctuations and on a European market, with very similar results, though no details are shown. **CRITIQUE** A sample of 406 stocks drawn from the 500 with a full history over 1991 to 1996

would exclude firms that were delisted, acquired or listed later, which is a survivorship filter that could raise measured correlations and change the market mode; the available text does not say how the 406 were chosen. No look-ahead problem arises because the exercise is descriptive. Data mining is limited by the absence of specification search, but the single free parameter σ^2 is chosen after inspecting the histogram.


4.3 Estimation and inference

There is no estimator in the usual sense. The spectrum is obtained by eigen-decomposition of C . The empirical density in the main figure is a smoothed histogram; the letter does not specify the smoothing, no standard errors or confidence bands are reported, and no formal goodness-of-fit statistic (for example a Kolmogorov-Smirnov distance between the empirical and theoretical spectral distributions) is computed. Inference is by visual agreement, which is acceptable as a diagnostic for a letter but is not a hypothesis test.

5. Results

5.1 The bulk of the spectrum follows the Marchenko-Pastur law

Tested. Whether the empirical eigenvalue density (Figure 1 of the letter) matches Equation (3). **Found.** The letter shows two theoretical curves: $\sigma^2 = 0.85$, obtained after removing the market mode (dotted curve), and a free fit. Letting σ^2 float, the best fit is $\sigma^2 = 0.74$ (solid curve), which the authors describe as accounting satisfactorily for 94 percent of the spectrum, leaving about 6 percent of eigenvalues above the upper edge. **Magnitude.** The largest eigenvalue is about 25 to 30 times the theoretical edge, roughly 61 in absolute terms, or about 15 percent of total variance in one direction (Section 3.3). **Economic meaning.** The market eigenvector alone therefore carries about $61/406 = 15$ percent of the total variance of the standardised system. The remaining eigenvalues of the bulk, from about 0.15 to about 1.8 under the fitted σ^2 , are the range one would see if the assets were unrelated apart from that factor. The fact that the whole low end lies inside the band means that no combination of these stocks was shown to have a genuinely low variance that pure noise could not also produce. **Statistical significance.** Not tested. **Support for the hypothesis.** Strong for the noise interpretation of the bulk, moderate for the claim of information above the edge, since no test of the excess is reported.


Figure 1. Simulated sample spectrum against the Marchenko-Pastur law, and the effect of Q 

Generated by the reviewer. Panel (a): eigenvalues of the sample correlation matrix of a simulated one-factor Gaussian panel ($N = 406$, $T = 1309$, single-factor correlation 0.148, seed 29092026), reproducing the size of the letter's sample and its reported $\sigma^2 = 0.85$. Panel (b): Equation (3) for $\sigma^2 = 1$. Not the letter's data.

Figure 1 is not the paper's figure. It is a controlled replication of the mechanism: a panel with one true factor and otherwise independent noise gives a largest eigenvalue of 60.4 and a bulk whose largest eigenvalue is 2.01, almost exactly the edge of 2.06 that Equation (4) predicts for $\sigma^2 = 0.85$. The simulated ratio of the top eigenvalue to the $\sigma^2 = 1$ edge is 24.9 and to the $\sigma^2 = 0.85$ edge 29.3, reproducing the letter's 25 and 30. This shows that the letter's headline numbers are what a single dominant factor plus noise would produce; it does not show that the true structure of the S&P 500 is a single factor, since the letter's bulk fit ($\sigma^2 = 0.74$) needs something more.

5.2 Eigenvectors inside the band look random; those above do not


Tested. Whether components of eigenvectors with eigenvalues in $[\lambda_{\min}, \lambda_{\max}]$ follow Equation (8). **Found.** The authors report that Equation (8) fits the empirical histogram of components extremely well for five eigenvectors chosen in the band, and that for the highest eigenvalues there are marked differences. **Meaning.** The loading vectors of the band eigenvalues are indistinguishable from random rotations, so the portfolios that these eigenvectors represent are not stable, economically meaningful combinations. The market eigenvector has roughly equal components on all stocks, so it is the average stock. **Robustness.** The test is graphical and uses five eigenvectors, not all of them.

Figure 2. Eigenvector components: inside the band and at the largest eigenvalue

Generated by the reviewer from the same simulated one-factor panel as Figure 1. Components are normalised so that the squared components sum to N . Panel (a): a bulk eigenvector against Equation (8). Panel (b): the eigenvector of the largest eigenvalue, whose components cluster near one, the equally weighted market direction. Not the letter's data.

5.3 What the paper reports about risk management

The paper does not compute portfolio risk. Its statement that the lowest eigenvalues, corresponding to the least risky Markowitz portfolios, are dominated by noise is an inference from the spectrum. The quantitative version is Equation (10), tested in Figure 3 and Table 2 with simulations that assume a true covariance equal to the identity.

Figure 3. Minimum-variance portfolio risk from an estimated matrix, in and out of sample

Generated by the reviewer. True covariance is the identity, $N = 200$, 60 replications per Q , variance measured relative to the true minimum $1/N$. Curves: Equation (10). Points: simulation. Not the letter's data.

Table 2. In-sample and out-of-sample variance of the estimated minimum-variance portfolio

Q = T/ N	In-sample (theory)	In-sample (simulated)	Out-of-sample (theory)	Out-of-sample (simulated)	Realised / predicted variance
1.25	0.200	0.207	5.000	4.969	25.0
1.50	0.333	0.331	3.000	2.975	9.0
2.00	0.500	0.505	2.000	1.980	4.0
3.22	0.690	0.702	1.450	1.443	2.1
5.00	0.800	0.806	1.250	1.259	1.6
10.00	0.900	0.898	1.111	1.111	1.2

Source: reviewer's calculation and simulation (N = 200, true covariance equal to identity, 60 replications). Variances relative to the true minimum. Last column uses theory: $1/(1 - 1/Q)^2$.

The result is stark at the Q of the letter's sample: a manager using the raw matrix would forecast a variance 31 percent below the best achievable and realise one 45 percent above it. This is estimation error under a noise-only truth; it is an upper bound on the damage for a market with a strong market factor, because the dominant factor is estimated precisely and the damage comes from the remaining $N - 1$ dimensions. The equal-weighted market direction itself is not the problem.

6. Robustness and Additional Tests

The letter reports few robustness checks, each mentioned rather than shown. The table sets out the concern, the test, the result, and whether it resolves the concern.

Table 3. Robustness tests in the letter

Concern	Test in the paper	Result	Resolves the concern?
Market mode contaminates the noise benchmark	Refit with $\sigma^2 = 1 - \lambda_1/N = 0.85$, then free $\sigma^2 = 0.74$	Fit improves; 94 percent of spectrum matched	Partly. Adds a free parameter and reveals that the null is rejected for the complement of the market
Results specific to price correlations	Same analysis on volatility correlation matrices from intraday data	Very similar results, no numbers shown	Not assessable without the evidence
Results specific to one market	Same analysis on another stock market (Paris mentioned)	Consistent, no numbers shown	Not assessable without the evidence
Eigenvectors in the band may contain information	Compare five in-band eigenvectors with Equation (8)	Good agreement except above $\lambda_{\max}$	Partly. Tests marginal distribution of components, not their structure

Concern	Test in the paper	Result	Resolves the concern?
Finite-size smoothing of the edges	Acknowledged that singularities are smoothed at finite N	No quantification	No. See Section 8 for a calculation showing the smoothing is negligible at N = 406
Time stability of the estimated correlations	Deferred as needing more careful study	None	No

Source: Laloux, Cizeau, Bouchaud and Potters (1999); assessment by the reviewer.

The main structural weakness of the robustness evidence is that every check is a repetition of the same graphical comparison on different data, none varies the null. Checks that would matter are alternative null models (heavy tails, volatility clustering, sector structure), alternative windows (subsamples of T), alternative universes (varying N at fixed T to trace the predicted change in Q), and a placebo, such as randomly permuting each stock's return series in time, which should reproduce the pure Marchenko-Pastur spectrum and would confirm that the fit is not an artefact of the smoothing or normalisation. Section 8 supplies simulated versions of some of these.

7. Critical Evaluation

7.1 What the paper actually proves

ESTABLISHED

In this sample, the bulk of the empirical eigenvalue distribution is well described by the Marchenko-Pastur density with a reduced variance parameter, and an estimated correlation matrix with $Q = 3.22$ has an eigenvalue spectrum whose lower and upper extremes are largely produced by sampling error. The market mode is far outside the noise band.

SUPPORTED BUT CONDITIONAL

The claim that about 6 percent of eigenvalues carry information depends on the null (independent, identically distributed, unit variance), on the visual fit, and on the choice of σ^2 . The claim that in-band eigenvectors are random depends on a marginal-distribution test on five eigenvectors.

SUGGESTIVE

The extension to volatility correlations and to other markets, since no evidence is presented. The idea that the same method could isolate meaningful correlations across asset classes, offered as a potential application.

NOT ESTABLISHED

That the raw correlation matrix produces worse portfolios than any alternative, since no portfolio is constructed. That eigenvalues above the edge are economically meaningful rather than caused by non-stationarity, heteroskedasticity or sector effects. That the result generalises to other sample sizes and periods. Any causal statement about the sources of correlation.

7.2 Strengths

Recognised by the authors. The parameter-light nature of the comparison, the agreement between theory and data over the whole bulk, and the potential to apply the tool to any type of assets.

Identified independently. (i) The question is sharply posed and answerable with one object, the spectrum. (ii) The benchmark is exact, not simulated, so it requires no computation and no free parameters other than σ^2 . (iii) The trace identity ties the fitted σ^2 to the variance share of the top eigenvalues, which lets the reader check the letter's own numbers for internal consistency, and they are consistent (Section 3.3). (iv) The eigenvector test is a second, independent diagnostic that would fail if the eigenvalue agreement were accidental. (v) The letter is short and every quantity needed to re-run it (N, T, Q, preprocessing) is reported, although the data are not.

7.3 Weaknesses

Most important weakness: the risk-management conclusion is not tested. *Problem:* the letter never builds a portfolio. *Why it matters:* the headline claim, that blind use of empirical matrices in risk management is unsafe, is inferred from the spectrum, and the size of the damage depends on how the optimiser uses the small eigenvalues, on constraints and on the true structure. *Potential bias:* the claim may overstate the damage for constrained (long-only) portfolios, which implicitly shrink weights, and understate it for unconstrained ones. *How it could be tested:* compute in-sample versus out-of-sample variance of minimum-variance portfolios on rolling windows with the raw matrix and a filtered one. Equation (10) and Figure 3 give the answer under an identity truth; the letter's data would give it for reality.

Weakness 2: the null model. *Problem:* the null assumes independent, identically distributed, homoskedastic returns. *Why it matters:* equity returns have volatility clustering and common volatility shocks, so a matrix of independent returns is a straw man. *Potential bias:* eigenvalues above the edge may reflect volatility co-movement rather than information about correlation. Section 8 shows a common stochastic volatility factor pushes the top eigenvalue to 5.3 with no correlation in levels. *Test:* rescale each day's returns by a cross-sectional volatility estimate before computing C.

Weakness 3: no statistical test and a free parameter. *Problem:* σ^2 is tuned by eye; the exceedance count (6 percent) is read from the plot. *Why it matters:* the number of information-bearing eigenvalues is the paper's main quantitative output. *Bias:* sensitivity to the fit (0.85 versus 0.74 gives different edges, 2.06 versus 1.79, and different counts). *Test:* a sequential test on the largest eigenvalue using the Tracy-Widom law (Johnstone, 2001) or the eigenvalue-ratio methods later developed.

Weakness 4: selection and non-stationarity. *Problem:* 406 survivors over a six-year window in one bull market. *Why it matters:* correlations rise in crises and drift between regimes; a single window cannot separate stable structure from regime structure. *Bias:* survivorship and period effects. *Test:* repeat over rolling and disjoint windows and compare eigenvector overlaps (Section 8).

Weakness 5: below-threshold information is invisible. *Problem:* the method labels as noise anything under the edge. *Why it matters:* real but weak factors are misclassified. *Bias:* the method is conservative about information. *Test:* spiked-model detection thresholds (Section 8).

Author-acknowledged limits are that the study of finite-size effects, stability over time, and volatility correlations requires more careful work. The weaknesses above are the reviewer's own, and none of the five is acknowledged in the letter.

8. Contradictions and Edge Cases

This section tries to break the result. All numbers are the reviewer's simulations ($N = 406$, $T = 1309$, seed 29092026) or calculations, with assumptions stated, and none is from the letter.

Table 4. Stress tests of the null-model interpretation

Case	Construction (reviewer's assumption)	Largest sample eigenvalue	Marchenko- Pastur edge ($\sigma^2 = 1$)	Reading
Pure noise, Gaussian	Independent standard normal	2.395	2.424	Inside the band, as theory says
Heavy tails, Student t with 3 degrees of freedom	Independent, scaled to unit variance	2.412	2.424	Inside the band in this single draw; infinite fourth moments could alter the edge in other samples
Common volatility factor	Independent returns scaled each day by $\exp(0.5 z_t)$, z_t standard normal	5.258	2.424	Well outside the band with no correlation in levels
Hidden sector, 20 stocks, intra- sector correlation 0.02	Population eigenvalue 1.38	2.402	2.424	Undetectable
Hidden sector, correlation 0.0293	Population eigenvalue 1.557, the detection threshold $1 + \sqrt{(N/T)}$	2.451	2.424	At the threshold, indistinguishable from the edge
Sector, correlation 0.05	Population eigenvalue 1.95	2.635	2.424	Detached from the band
Sector, correlation 0.10	Population eigenvalue 2.90	3.287	2.424	Clearly detached

Source: reviewer's simulation. Detection threshold from the spiked-covariance results of Baik, Ben Arous and P\'ech\'e (2005).

Under what conditions should the result disappear? The result is a statement about Q . As $Q \rightarrow 1$ the lower edge goes to zero and the whole density gets pushed against zero (Figure 1, panel b); for $Q < 1$ the matrix is singular and $N - T$ eigenvalues are exactly zero, so the letter's framework says nothing about the level of the smallest eigenvalues when $N > T$. As Q grows the band narrows to a point at σ^2 and the noise interpretation disappears, which is why the paper's conclusion is specific to a sample with Q of about three.

Could another mechanism generate the same result? Yes. Table 4 shows that a common stochastic volatility factor with no correlation in levels produces a top eigenvalue more than twice the noise edge. The letter's above-edge eigenvalues are therefore compatible with correlation in levels, correlation in volatilities, or both, and the design cannot tell them apart. A second alternative: the bulk misfit that the authors attribute to $\sigma^2 = 0.74$ might arise from a mildly heterogeneous population spectrum (many weak sectors) rather than pure noise; the Marchenko-Pastur shape is then a misspecified null, and the fitted σ^2 has no clean interpretation.

Weak but genuine factors are invisible. The phase-transition threshold for detection in the spiked model is a population eigenvalue of $1 + \sqrt{(N/T)} = 1.557$. A sector of 20 stocks would need an average intra-sector correlation of $(1.557 - 1)/19 = 0.029$ to reach it, so any sector-wide correlation below about 0.03 is unrecoverable at this Q, and a correlation of 0.05 produces a detached eigenvalue of only 2.63 (Table 4). Anything below the edge is classed as noise by construction, so the method is conservative about information, and the count of information-bearing eigenvalues (6 percent) is a lower bound on the true factor number.

Finite-size edge smoothing is negligible. The letter notes that finite-N smoothing of the edge exists but does not quantify it. The fluctuation scale of the largest eigenvalue under the Gaussian null is of order $N^{-2/3}$; applying Johnstone's centring and scaling (calculation) gives a centring of 2.423 and a scale of about 0.018 for $N = 406$, $T = 1309$. An eigenvalue must therefore exceed the edge by only a few hundredths to be outside the band, so smoothing cannot explain the letter's 24 or so excess eigenvalues, although the counting is sensitive to the fitted σ^2 (Section 7.3).

Are the in-band eigenvectors really unstable? In the simulated one-factor panel, splitting the sample in two, the overlap (absolute inner product) of the largest eigenvector across halves is 0.994, while for the eigenvectors ranked 101 and 406 it is 0.007 and 0.074, consistent with the $1/\sqrt{N} \approx 0.05$ of random directions. This supports the interpretation of Equation (8) and adds the point the letter defers: noise eigenvectors are not stable over time. Whether this holds in the real market cannot be established from the letter.

Does it extend outside the setting? The method is valid for any panel with a well-defined N and T , but the independence null is more damaging in markets with strong common volatility, and Q is small in exactly the settings, such as hundreds of assets with monthly data, where it matters most.

9. Further Literature

Each row gives the paper, the idea it shares with the letter, its relationship to it, and why the connection matters. Only connections with a substantive link are included.

Table 5. Literature connections

Paper	Shared idea	Relationship	Why it matters
Marchenko and Pastur (1967)	Limiting eigenvalue law of sample covariance matrices	Mathematical foundation of Equations (3) and (4)	Supplies the exact noise benchmark; the letter contributes its transfer to finance
Markowitz (1952)	Mean-variance selection using the covariance matrix	Motivation	The optimiser's sensitivity to small eigenvalues is the mechanism behind the letter's warning
Elton and Gruber (1973)	Structured estimators of the covariance matrix for portfolios	Precursor	Earlier evidence that unstructured sample matrices are poor inputs; the letter offers a diagnostic rather than a model
Jobson and Korkie (1980); Michaud (1989)	Estimation error in optimised portfolios	Precursor, empirical	Document the consequence of noise, that the letter explains through the spectrum
Chamberlain and Rothschild (1983)	Approximate factor structure with a few large eigenvalues	Theoretical counterpart	A few eigenvalues that diverge with N while the rest are bounded is the population analogue of the letter's picture
Plerou et al. (1999)	Random matrix analysis of cross-correlations	Simultaneous, independent	Same finding on different data in the same issue; adds eigenvector localisation and stability
Laloux et al. (2000)	Same authors, longer treatment with a cleaning procedure	Direct extension	Supplies the portfolio-risk tests the letter lacks
Plerou et al. (2002)	Extensive random matrix analysis of US equity correlations	Extension	Tests deviations from the null in detail, including sector and time structure
Johnstone (2001); Baik, Ben Arous and P��ch�� (2005)	Distribution and phase transition of largest eigenvalues	Refines the theory	Give the finite-N edge and the detection threshold used in Section 8
Ledoit and Wolf (2004)	Shrinkage covariance estimator	Alternative methodology	Improves portfolio risk without the spectrum decomposition; competitor to eigenvalue cleaning
Pafka and Kondor (2004)	Portfolio risk under estimated correlation matrices	Quantifies the letter's implication	Source of the in-sample and out-of-sample results in Equation (10)
Bun, Bouchaud and Potters (2017)	Rotationally invariant cleaning of correlation matrices	Later development by the same school	Shows optimal eigenvalue cleaning is a shrinkage; replaces the clipping that the letter suggests

Paper	Shared idea	Relationship	Why it matters
DeMiguel, Garlappi and Uppal (2009)	Optimal versus naive diversification	Contradictory or sobering evidence	Estimation error is so large that equal weighting often beats optimisation, which limits what any covariance cleaning can achieve

Source: reviewer's assessment. Full references in the bibliography.

10. Research Gaps

10.1 Paper-level gaps

- **Gap: no portfolio test.** Why it matters: the practical claim is unproven. Existing limitation: no in-sample versus out-of-sample comparison. Way forward: rolling-window minimum-variance tests with and without cleaning.
- **Gap: no statistical test of the null.** Why it matters: the count of informative eigenvalues drives the message. Existing limitation: visual fit with a tuned σ^2 . Way forward: Tracy-Widom or eigenvalue-ratio tests with volatility-standardised returns.
- **Gap: time variation.** Why it matters: correlations are non-stationary, and the noise band depends on the window. Existing limitation: one window. Way forward: rolling spectra and eigenvector overlaps.

10.2 Literature-level gaps

- **Gap: separating correlation from co-volatility.** Why it matters: eigenvalues above the edge can arise from common volatility (Table 4). Existing limitation: most analyses use raw returns. Way forward: joint models of level and volatility structure.
- **Gap: economic content of the sub-leading eigenvalues.** Why it matters: the non-market eigenvalues in the letter are only modestly above the edge (Section 3.3). Existing limitation: sector labels are used post hoc. Way forward: link eigenvectors to observed fundamentals with out-of-sample validation.

10.3 Methodological gaps

- **Gap: non-Gaussian and dependent data.** Why it matters: the Marchenko-Pastur law needs independence. Existing limitation: asymptotics for dependent, heavy-tailed data are harder. Way forward: simulation-calibrated nulls that preserve the marginal and volatility dynamics of each stock, for example by block-permuting time.
- **Gap: unbalanced and asynchronous data.** Why it matters: missing observations and asynchronous trading reduce the effective T. Existing limitation: the letter assumes a full panel. Way forward: high-frequency realised covariances with Epps-effect corrections.

11. Research Opportunities

Five ideas that emerge from the analysis are set out below; each states the data and identification used. Data sources are named as candidates, not as retrieved or verified datasets, and no result is claimed for any of them.

Idea 1. Out-of-sample value of spectral cleaning versus shrinkage

RESEARCH QUESTION. Does removing eigenvalues inside the Marchenko-Pastur band reduce the realised variance of minimum-variance portfolios relative to the raw matrix and to Ledoit-Wolf shrinkage, as a function of Q and of the market regime?

MOTIVATION. The letter implies a benefit but never measures it; later papers propose competing cleaners.

HYPOTHESIS. Cleaning reduces realised variance most where Q is below about 3, and the gain shrinks as Q grows, with shrinkage matching cleaning when the population spectrum has no spikes.

ECONOMIC MECHANISM. Estimation error in the smallest eigenvalues is amplified by the $1/\lambda$ weighting in Equation (9).

LITERATURE CONNECTION. Laloux et al. (2000); Ledoit and Wolf (2004); Bun, Bouchaud and Potters (2017); DeMiguel et al. (2009).

DATA. CRSP daily stock returns, 1990 to 2025, candidate source, US common stocks, requiring access.

VARIABLES. Realised out-of-sample portfolio variance, predicted variance, turnover, Q .

METHOD. Rolling estimation windows, N drawn at random from the universe at several sizes, several estimators.

IDENTIFICATION STRATEGY. Random draws of stocks and windows give repeated, comparable trials; pairwise comparisons within trials remove universe effects.

EXPECTED CONTRIBUTION. A calibrated map of when spectral cleaning is worth using.

MAIN RISK. Crowded literature; gains may vanish once constraints and transaction costs are included.

Idea 2. A volatility-robust test for the number of information-bearing eigenvalues

RESEARCH QUESTION. How many eigenvalues of the US equity correlation matrix lie outside the noise band once common volatility movements are filtered out?

MOTIVATION. Table 4 shows that a common volatility factor alone produces an eigenvalue more than twice the edge; the letter's count of about 24 has never been shown to be free of this.

HYPOTHESIS. After volatility standardisation the count falls, and the fall is largest in high-volatility periods.

ECONOMIC MECHANISM. Volatility co-movement, driven by common shocks to risk, inflates sample correlations without altering true correlations in levels.

LITERATURE CONNECTION. Plerou et al. (2002); Johnstone (2001); Bai and Ng (2002) on the number of factors.

DATA. Daily returns from CRSP or a comparable public source, 1991 to 2025, candidate source; realised volatility from intraday data if available.

VARIABLES. Eigenvalue counts, top-eigenvalue share, common-volatility proxy.

METHOD. Standardise returns by rolling, cross-sectional volatility estimates; apply sequential Tracy-Widom tests; compare with the raw-return counts.

IDENTIFICATION STRATEGY. The difference in counts between raw and standardised returns, across volatility regimes, isolates the co-volatility contribution.

EXPECTED CONTRIBUTION. A defensible estimate of the noise band and factor count, which could be reused by risk models.

MAIN RISK. The standardisation itself may remove genuine correlation dynamics.

Idea 3. Eigenvector persistence as a measure of information

RESEARCH QUESTION. Do eigenvectors outside the noise band persist across adjacent windows, and does persistence predict out-of-sample covariance forecasting accuracy?

MOTIVATION. The letter tests eigenvectors only within one window; persistence provides an out-of-sample criterion for information.

HYPOTHESIS. Outside-band eigenvectors have high overlap across windows, band eigenvectors have overlap of order $1/\sqrt{N}$, and the number of persistent eigenvectors exceeds the count predicted by the edge alone.

ECONOMIC MECHANISM. Persistent factors arise from stable economic exposures; noise eigenvectors change with the sample.

LITERATURE CONNECTION. Plerou et al. (2002); Bun, Bouchaud and Potters (2017).

DATA. Same as Idea 1.

VARIABLES. Overlap between eigenvectors from adjacent windows; forecast error of realised covariance.

METHOD. Rolling windows, matched eigenvector pairs, regression of forecast error on overlap.

IDENTIFICATION STRATEGY. Comparison of persistence inside and outside the band within the same windows.

EXPECTED CONTRIBUTION. A data-driven definition of informative eigenvectors.

MAIN RISK. Eigenvalue crossings and rotations in degenerate subspaces make matching ambiguous.

Idea 4. The market-mode share as a regime indicator

RESEARCH QUESTION. Does the share of variance in the largest eigenvalue, λ_1/N , predict subsequent diversification failure and drawdowns, across countries and asset classes?

MOTIVATION. The letter's $\lambda_1/N \approx 0.15$ is a single point; its time variation is a natural state variable.

HYPOTHESIS. Rises in λ_1/N precede higher realised portfolio volatility beyond that implied by index volatility.

ECONOMIC MECHANISM. Increased common exposure reduces the effective number of independent bets.

LITERATURE CONNECTION. Plerou et al. (1999); Chamberlain and Rothschild (1983).

DATA. Index constituents for several countries from public index providers, candidate source; requires access.

VARIABLES. λ_1/N , second to fifth eigenvalue shares, forward realised volatility, drawdown.

METHOD. Predictive regressions with overlapping-observation standard errors, out-of-sample tests.

IDENTIFICATION STRATEGY. Cross-country and cross-period variation, with volatility controls to separate it from index volatility.

EXPECTED CONTRIBUTION. A model-free indicator of correlation risk.

MAIN RISK. Mechanically related to market volatility; incremental information may be nil.

Idea 5. Larger Q from intraday data: how far does the noise band shrink?

RESEARCH QUESTION. When realised covariances from intraday data replace daily ones, so that the effective Q is much larger, how many eigenvalues remain outside the band, and how do their eigenvectors compare with the daily ones?

MOTIVATION. The noise band narrows as Q grows (Figure 1b), so intraday data should reveal structure that daily data hide.

HYPOTHESIS. The number of outside-band eigenvalues increases with sampling frequency until microstructure noise and asynchronicity dominate.

ECONOMIC MECHANISM. More observations per asset reduce sampling error, but asynchronous trading biases correlations toward zero.

LITERATURE CONNECTION. Laloux et al. (2000); Plerou et al. (2002); the letter's own reference to volatility correlations from intraday fluctuations.

DATA. Trade-and-quote data from a commercial vendor (for example TAQ), candidate source; requires access.

VARIABLES. Eigenvalue counts by sampling frequency, overlap with daily eigenvectors.

METHOD. Realised covariance at 5-, 15- and 30-minute grids with asynchronicity corrections.

IDENTIFICATION STRATEGY. Within-sample variation across sampling frequencies, holding the universe fixed.

EXPECTED CONTRIBUTION. Establishes an upper bound on the number of recoverable factors.

MAIN RISK. Microstructure effects can mimic or hide structure.

11.1 Ranking

Table 6. Ranking of research ideas (1 = weak, 5 = strong; reviewer's judgement, not measurement)

Idea	Novelty	Feasibility	Data availability	Method quality	Economic importance	Contribution	Total
1. Cleaning versus shrinkage	2	5	4	4	4	3	22
2. Volatility-robust factor count	4	4	4	4	4	4	24
3. Eigenvector persistence	4	4	4	3	3	4	22
4. Market-mode share as indicator	3	5	3	3	4	3	21
5. Intraday data	3	2	2	4	3	4	18

Source: reviewer's scoring; totals are simple sums.

The top three are Idea 2, Idea 1 and Idea 3 (Idea 1 and Idea 3 tie at 22; Idea 1 is placed second because it is more feasible). Idea 2 ranks first because it attacks the largest unresolved weakness of the letter (the null model) with public data and existing tools, and its result would feed directly

into every other idea. Idea 1 is the most realistic quantitative-finance project and produces a directly usable result. Idea 3 gives a non-parametric definition of information that avoids the null-model debate.

12. Learning Framework

12.1 Prerequisite concepts

Linear algebra (eigen-decomposition of symmetric matrices, trace, orthogonal rotation); sample covariance and correlation estimation; Markowitz mean-variance optimisation; the Wishart distribution; basic probability on limits of empirical distributions; entropy maximisation.

12.2 Core concepts learned

Sampling noise creates a deterministic eigenvalue band (the Marchenko-Pastur band) whose width is set by Q ; the smallest and largest sample eigenvalues are biased outward; noise eigenvectors are random rotations; the trace identity ties the variance of the bulk to that of the outliers; the small eigenvalues control minimum-variance portfolios and are the most contaminated.

12.3 Important equations

- Equation (3) with Equation (4): the noise density and its edges.
- Equation (6): $\sigma^2 = 1 - \lambda_1/N$, the variance left for the bulk.
- Equation (8): the Porter-Thomas law for random eigenvector components.
- Equations (9) and (10): the minimum-variance solution and the in-sample versus out-of-sample gap.

12.4 Methods learned

Spectral comparison with a random-matrix null; fitting the bulk with a scale parameter; graphical eigenvector testing; using an exact asymptotic law as a benchmark rather than simulating.

12.5 Questions to test understanding

COMPREHENSION

- Why do the eigenvalues of a correlation matrix estimated from independent data spread over an interval rather than concentrate at one?
- What is the meaning of Q , and what happens to the spectrum as Q approaches one?
- Why does the letter report both 25 and 30 as the ratio of the top eigenvalue to the edge?

ANALYSIS

- Derive the edges in Equation (4) from the discriminant of Equation (5).

- Why does the trace identity imply that the bulk fitted with $\sigma^2 = 0.74$ holds 74 percent of variance?
- Why does a minimum-variance optimiser over-weight the smallest eigenvalue directions, and how does Equation (10) quantify the effect?

RESEARCH

- Which alternative null would you use for returns with volatility clustering, and how would you calibrate it?
- How would you design a test showing that an eigenvalue above the edge reflects correlation and not co-volatility?
- What would the letter's analysis look like for N greater than T ?

12.6 Papers to read next

Laloux et al. (2000); Plerou et al. (2002); Ledoit and Wolf (2004); Pafka and Kondor (2004); Bun, Bouchaud and Potters (2017); Johnstone (2001).

13. Conclusion

One-sentence thesis

The sample correlation matrix of a few hundred stocks over a few years is, apart from a market mode and a small number of other eigenvalues, statistically indistinguishable from the correlation matrix of independent noise, and the low-risk end of its spectrum is contaminated in particular.

Three most important insights

- Noise has a shape: the band $[\lambda_{\min}, \lambda_{\max}]$ is fixed by Q alone, so it can be computed rather than simulated.
- Most eigenvalue structure that looks like information is inside the band and is not stable.
- The risk-management implication rests on Equation (9): the optimiser weights directions by $1/\lambda$, and the smallest sample eigenvalues are the least reliable.

Strongest evidence

The agreement between the empirical density and the Marchenko-Pastur curve over 94 percent of the spectrum, together with the trace-identity consistency between $\sigma^2 = 0.74$ and the 26 percent of variance in the top 6 percent of eigenvalues, and the Gaussian shape of in-band eigenvector components.

Biggest weakness

The letter never tests its risk-management conclusion on portfolios, and its null model omits volatility clustering, which can itself produce eigenvalues above the edge (Table 4).

Most important research gap

A statistically valid count of information-bearing eigenvalues, robust to common volatility and to non-stationarity.

Best research opportunity

Idea 2 in Section 11: a volatility-robust test of the factor count on public equity data.

Three papers to read next

Laloux et al. (2000), Plerou et al. (2002), and Bun, Bouchaud and Potters (2017).

14. References

-
- Bai, J., and Ng, S. (2002). Determining the number of factors in approximate factor models. *Econometrica*, 70(1), 191–221.
- Baik, J., Ben Arous, G., and Péché, S. (2005). Phase transition of the largest eigenvalue for non-null complex sample covariance matrices. *Annals of Probability*, 33(5), 1643–1697.
- Bun, J., Bouchaud, J.-P., and Potters, M. (2017). Cleaning large correlation matrices: tools from random matrix theory. *Physics Reports*, 666, 1–109.
- Chamberlain, G., and Rothschild, M. (1983). Arbitrage, factor structure, and mean-variance analysis on large asset markets. *Econometrica*, 51(5), 1281–1304.
- DeMiguel, V., Garlappi, L., and Uppal, R. (2009). Optimal versus naive diversification: how inefficient is the 1/N portfolio strategy? *Review of Financial Studies*, 22(5), 1915–1953.
- Elton, E. J., and Gruber, M. J. (1973). Estimating the dependence structure of share prices: implications for portfolio selection. *Journal of Finance*, 28(5), 1203–1232.
- Jobson, J. D., and Korkie, B. (1980). Estimation for Markowitz efficient portfolios. *Journal of the American Statistical Association*, 75(371), 544–554.
- Johnstone, I. M. (2001). On the distribution of the largest eigenvalue in principal components analysis. *Annals of Statistics*, 29(2), 295–327.
- Laloux, L., Cizeau, P., Bouchaud, J.-P., and Potters, M. (1999). Noise dressing of financial correlation matrices. *Physical Review Letters*, 83(7), 1467–1470. DOI 10.1103/PhysRevLett.83.1467. Preprint arXiv:cond-mat/9810255.
- Laloux, L., Cizeau, P., Potters, M., and Bouchaud, J.-P. (2000). Random matrix theory and financial correlations. *International Journal of Theoretical and Applied Finance*, 3(3), 391–397.

- Ledoit, O., and Wolf, M. (2004). A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis*, 88(2), 365–411.
- Marchenko, V. A., and Pastur, L. A. (1967). Distribution of eigenvalues for some sets of random matrices. *Matematicheskii Sbornik*, 72(114), 507–536.
- Markowitz, H. (1952). Portfolio selection. *Journal of Finance*, 7(1), 77–91.
- Michaud, R. O. (1989). The Markowitz optimization enigma: is optimized optimal? *Financial Analysts Journal*, 45(1), 31–42.
- Pafka, S., and Kondor, I. (2004). Estimated correlation matrices and portfolio optimization. *Physica A*, 343, 623–634.
- Plerou, V., Gopikrishnan, P., Rosenow, B., Amaral, L. A. N., and Stanley, H. E. (1999). Universal and nonuniversal properties of cross correlations in financial time series. *Physical Review Letters*, 83(7), 1471–1474.
- Plerou, V., Gopikrishnan, P., Rosenow, B., Amaral, L. A. N., Guhr, T., and Stanley, H. E. (2002). Random matrix approach to cross correlations in financial data. *Physical Review E*, 65, 066126.
- Bibliographic note. The letter was accessed in its arXiv form. Details of the other references are given from the reviewer's knowledge of the literature and were not re-verified against the publishers' pages in this session, except that the letter's own journal reference and DOI were confirmed from its arXiv listing.*